

ECEN 449 – Microprocessor System Design



FPGAs and Reconfigurable Computing

Some of the notes for this course were developed using the course notes for ECE 412 from the University of Illinois, Urbana-Champaign

ROADMAP

I

now slide #s

1-18 : how to choose the right platform for a digital system

I

later

19-31 : under the hood of the FPGA
(ckt level)

II

32-end : CAD flow for FPGAs
("what goes on inside Vivado")

III

- ✓ FPGA = Field Programmable Gate Array
- ✓ ASIC = Application Specific Integrated Circuits
(automated design process)
- ✓ Custom IC = more manual than ASIC
- ✓ NRE = non-recurring engineering expense

(1)

which platform to use?

Criteria	Software	FPGA	ASIC	Custom IC
Delay	100+	10 - 15	2	1
time to design	1	2	10	30
power	100	10	1.5	1
area	<u>1000+</u>	10	1.5	1
design team size	1	2	10	50
cost = $\alpha B + \gamma$	2	9	120/200	1600/2500
NRE	1	5	20/100	100/1000
ease of upgrade-ability	1000	100	1	1
recurring cost	1	2	10	50

fabless
 • AMD
 • Nvidia
 • Apple

fab-based
 • intel
 • Samsung

Objectives of this Lecture Unit

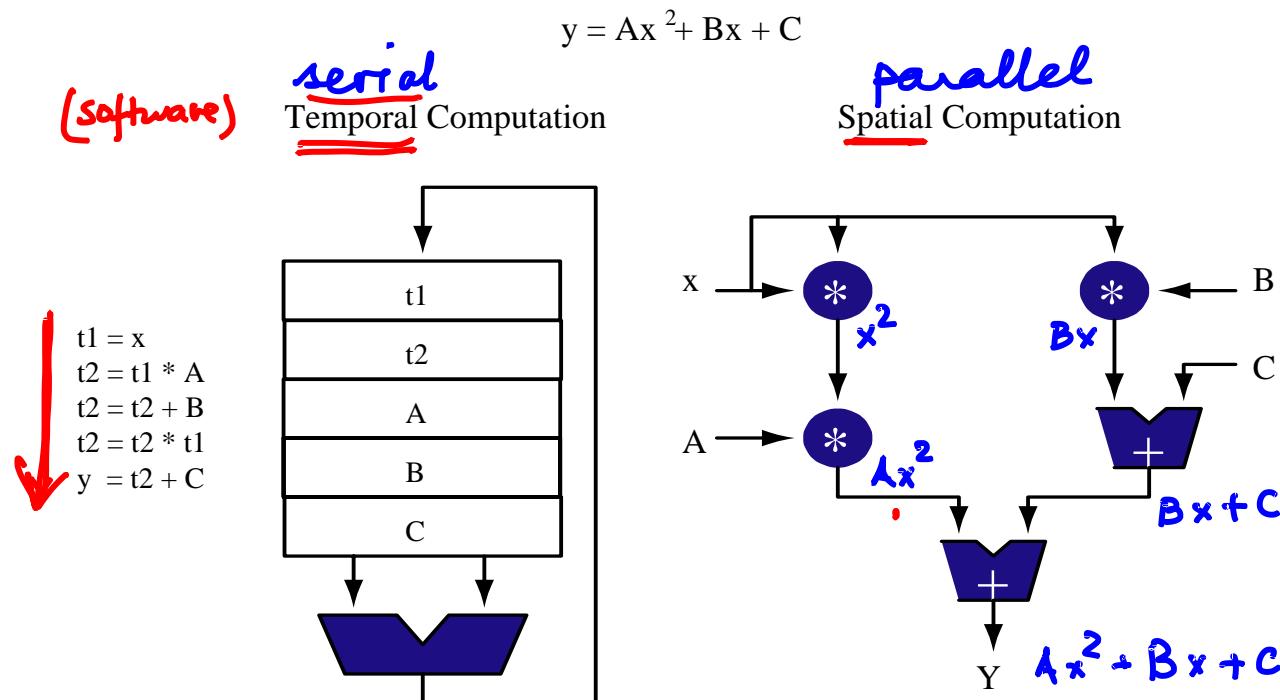
- Get a feel for the different technologies that can be used to implement a **digital** design (**system**):
 - Flavors of hardware technologies
 - Flavors of implementation methods
- Understand the basics of how FPGAs work
 - So that the CAD tools in the lab make sense to you

Software, Custom Hardware or Reconfigurable Hardware?

- When should we use software, “custom” hardware, or reconfigurable hardware?
- Software based systems are easiest to implement
 - But there is a huge performance gap between software and hand-designed (custom) hardware systems
 - Often 100-to-1 ratio of performance (speed) or performance/area
- But custom hardware systems not so good for general computing
 - Big design effort (time, cost) are barriers to implementation
 - Not practical to buy a new machine every time you want to run a different program
- Reconfigurable systems offer best-of-both-worlds
 - Run-time programmability (in the field)
 - Hardware-level performance (although lower than custom hardware)
 - FPGAs and CPLDs are the vehicles for reconfigurable systems.

Why is Hardware Faster Than Software?

- Spatial vs. Temporal Computation
 - Processors divide computation across time, dedicated hardware divides across space
 - But dedicated hardware is hardwired for a specific task.



Why is Hardware Faster Than Software?

- Specialization:
 - Instruction set may not provide the operations your program needs
 - Processors provide hardware that may not be useful in every program or in every cycle of a given program
 - Multipliers
 - Dividers
- Instruction Memory
 - Processors need lots of memory to hold the instructions that make up a program and to hold intermediate results.
- Bit Width Mismatches
 - In general, processors have a fixed bit width, and all computations are performed on that many bits
 - Multimedia vector instructions (MMX) a response to this

So why not just use Hardware?

- Dedicated hardware is
 - Dedicated (not flexible)
 - Takes long to design and develop (typical processor takes a handful of years to design, with design teams of a few hundred engineers)
 - This is expensive!
 - Only way to justify such an effort is if the customer demand guarantees high volume sales
- So there is a strong need for a design approach which has performance comparable to dedicated hardware, with ease-of-programmability comparable to software. **PLA**
- Answer? Reconfigurable computing (FPGAs, CPLDs and their cousins)
*complex
not logic devices*

Good Applications for Reconfigurable Computing

- Data Parallelism
 - Execute same computations on many independent data elements
 - Pipeline computations through the hardware
- Small and/or varying bit widths
 - Take advantage of the ability to customize the size of operators
- Low-volume applications which require rapid design turn-around time and hardware-like speeds
 - Several telecom, DSP (filters), radar, genomics (DNA sequence matching), processor emulation, neural network and similar applications.

Will FPGAs Defeat CPUs?

- Capacity: Instructions are very dense representation, logic blocks aren't
- Tools: Compilers for reconfigurable logic aren't very good
 - Some operations are hard to implement on FPGAs
 - C-for-FPGA technology is improving fast, though

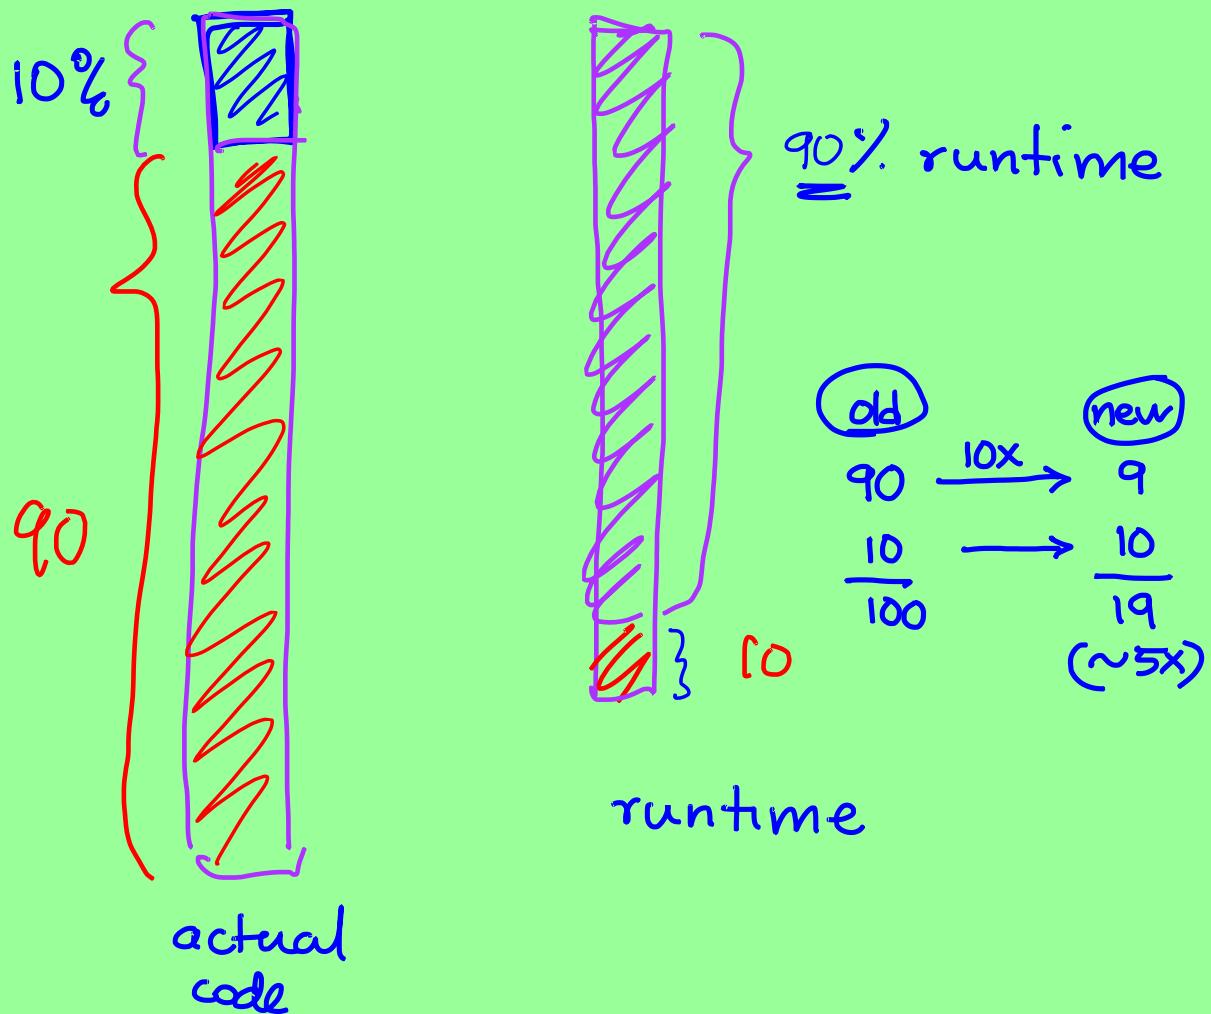
One approach to capacity is to exploit the ~~90-10 rule~~ of software

- Run the 90% of code that takes 10% of execution time on a conventional processor
- Run the 10% of code that takes 90% of execution time on reconfigurable logic
- But the temptation to merge the two worlds is real
 - Programmable-reconfigurable processors

90-10 rule

8.1

Digital system in software



8.2

Amdahl's law

Serial code = S time units
(red)

Parallelizable code (blue) = P time units

Old runtime was $S+P$

Suppose we accelerate the parallelizable part by a factor F

$$\text{New runtime} = S + \frac{P}{F}$$

8.3

speedup \leftrightarrow

$$\frac{S+P}{S+P/F} = \frac{\text{old runtime}}{\text{new runtime}}$$

for the 90-10 example

(F = 10 assumed)

$$\frac{\text{red}}{\text{old}} \frac{10 + 90}{10 + 90/10} = \frac{100}{10}$$

~5

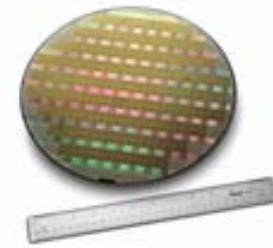
$$\text{max speedup} = \frac{100}{10 + 90/10} \approx 10$$

A Peek Under the Hood

- In the next few slides, we will peek under the hood of some of competing hardware based digital system design platforms
- We will cover
 - Application Specific ICs (ASICs).
 - Examples are IP routing ICs
 - SSI/MSI/LSI/VLSI
 - Reconfigurable (also sometimes called programmable) ICs.
 - Examples are FPGAs, CPLDs
 - Full custom Integrated Circuits (ICs).
 - Examples are processors, GPUs, network processors, DSP processors.

Application Specific Integrated Circuits

- Very high capacity today -- 10-100M transistors
- Very high speed – 500MHz+
 - Integration
 - Specificity
- Can use any design style below (or a hybrid)
 - Full Custom
 - Standard-cell (synthesized) – dominating methodology due to manufacturing considerations
- Long fabrication time
 - Weeks-months from completed design to product
- Only economical for high-volume parts
 - Making the masks required for fabrication is becoming very expensive, in the order of \$1M per design

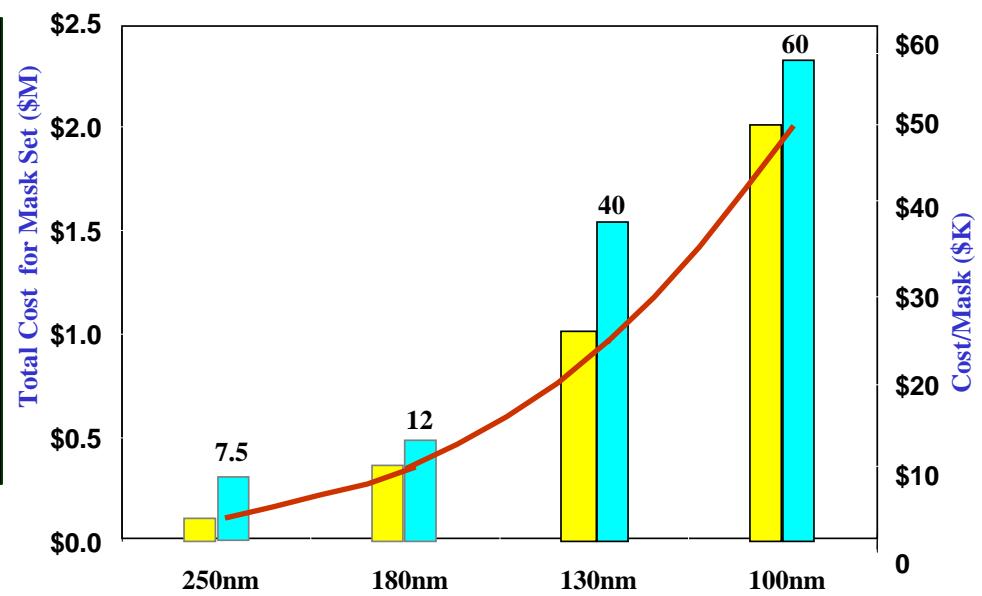


Deep Submicron Design Challenges

- This slide discusses why ASICs are becoming less popular in recent times (compared to reconfigurable ICs)
- Physical effects are increasingly significant
 - Parasitics, reliability issues, power management, process variation, etc.
- Design complexity is high
 - Multi-functionality integration
 - Design verification is a major limitation on time-to-market
- Cost of fabrication facilities and mask making has increased significantly

Rapid Increase in Manufacturing Cost

Process (um)	2.0	...	0.8	0.6	0.35	0.25	0.18	0.13	0.10
Single Mask cost (\$K)	1.5		1.5	2.5	4.5	7.5	12	40	60
# of Masks	12		12	12	16	20	26	30	34
Mask Set cost (\$K)	18		18	30	72	150	312	1,000	2,000

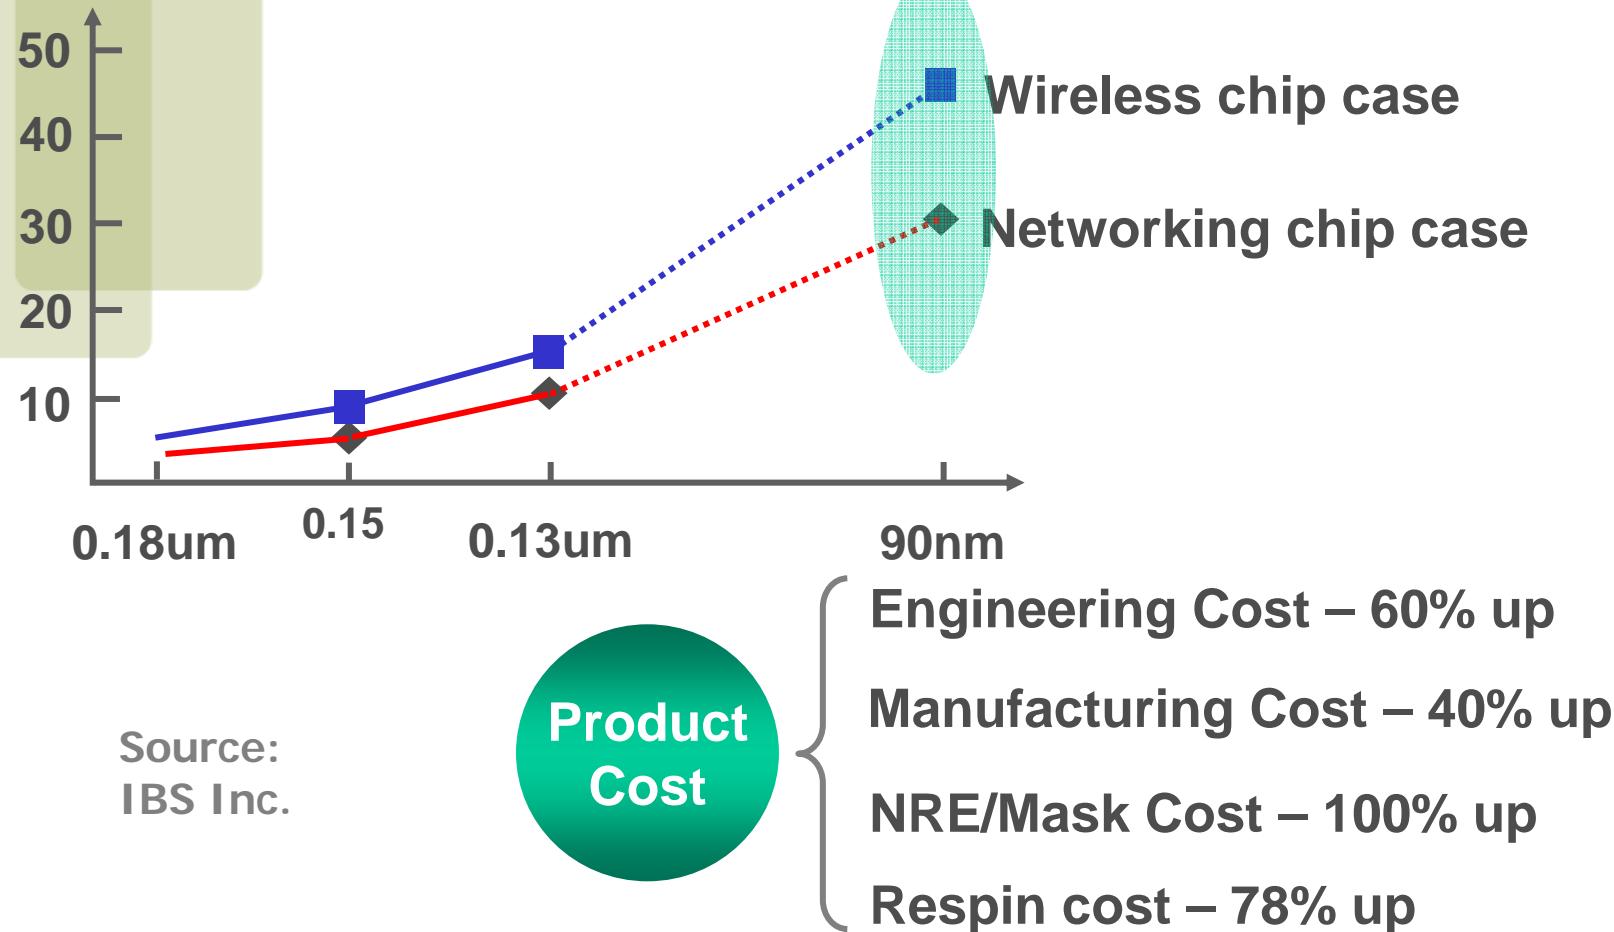


Source: EETimes

The Cost of Next Generation Product

Total Product Cost (\$M)

\$30M ~ \$50M @ 90nm



Programmable Logic Devices

- Early version: Mask-Programmable Gate Arrays
 - Build standard layout of transistors on chip
 - Customer specifies wiring to connect transistors into gates/system
 - Only has to go through last few mask steps of fabrication process, so faster than full chip fabrication
 - May become popular again in the near future
- Newer version: Programmable Logic Devices (PLD) **PLA [248]**
 - Use AND-OR array to implement arbitrary Boolean functions
 - Programmed by burning fuses that define connection from input wires to gates
 - Customer site programming allows rapid prototyping
 - Limited capacity, functionality
 - Generally have to be used in conjunction with other parts to hold state
 - Used to implement logic with moderate number of inputs (< 20)

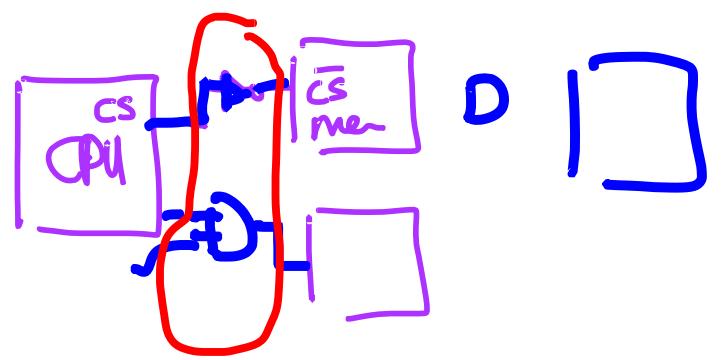
Programmable Logic Device Advantages



- Short TAT (total turnaround time)
- No or very low NRE (non-recurring engineering) costs.
- Field-reprogrammable
- Platform-based design

Today - Two Major Types of Programmable Logic

- CPLD (complex programmable logic device)
 - coarse-grained two-level AND-OR programmable logic arrays (PLAs)
 - fast and more predictable delay
 - simpler interconnect structures
- FPGA (field programmable gate array)
 - fine-grained logic cells
 - high logic density
 - good design flexibility (field programmable), easy redesign (just reprogram the chip!)
 - arguably more popular
- *Increasing ASIC design costs are making FPGAs more popular. This technology is therefore important to learn about. Hence this course.*
 - *Enables “garage” technology companies to thrive. This has a huge impact.*



Evolution of the FPGA

- Early FPGAs used mainly for “glue logic” between other components
 - Simple CLBs, small number of inputs
 - Focus was on implementing “random” logic efficiently
- As capacities grew, other applications emerged
 - FPGAs as alternative to custom IC’s for entire applications
 - Computing with FPGAs
- FPGAs have changed to meet new application demands
 - Carry chains, better support for multi-bit operations
 - Integrated memories, such as the block RAMs in the devices we’ll use
 - Specialized units, such as multipliers, to implement functions that are slow/inefficient in CLBs
 - *Newer devices incorporate entire CPUs: Xilinx Virtex II Pro has 14 ARM PowerPC CPUs (we will use such a device in our lab!!!)*

Full Custom ICs

- These have captured an important niche in hardware implementation of systems
- Microprocessors, GPUs, network processors, DSP processors are key examples.
 - HIGH sale volume (required to justify huge development cost and time)
 - These are often the flagship products of many semiconductor companies (Intel, IBM, AMD, TI, Freescale, etc)
 - These designs are “custom” designed, to do a specific task extremely fast, with minimum area and power.

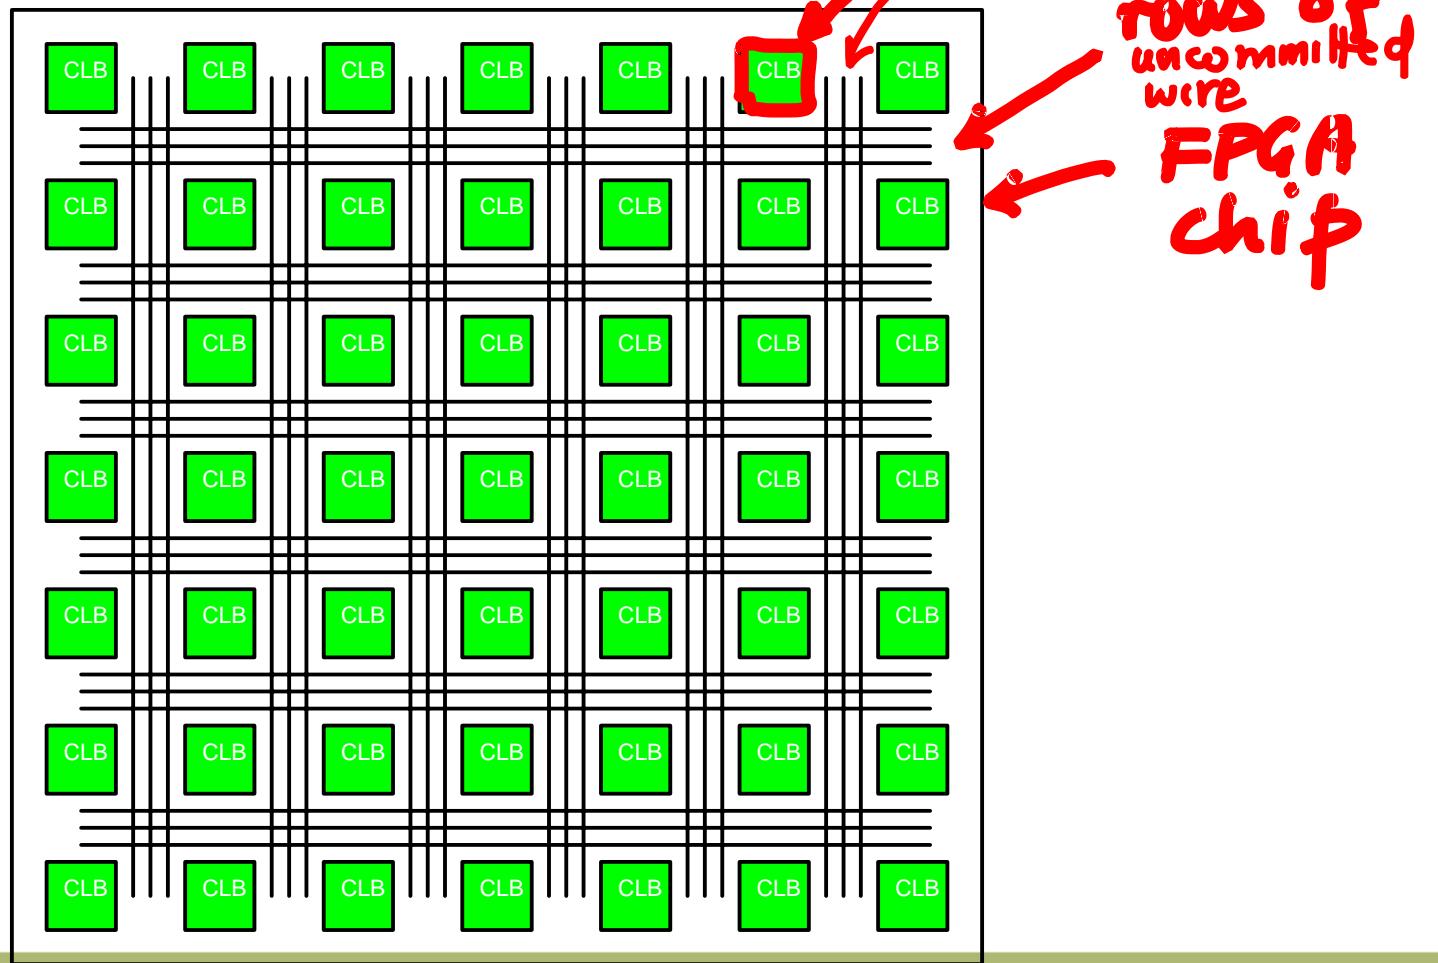
FPGAs in Detail

Inside

- Now in the next few slides, we will look at the technology that is inside an FPGA IC.
- This will allow us to understand how the FPGA works
- After this, we will be able to make sense of the design flow that is used to design a FPGA based circuit.

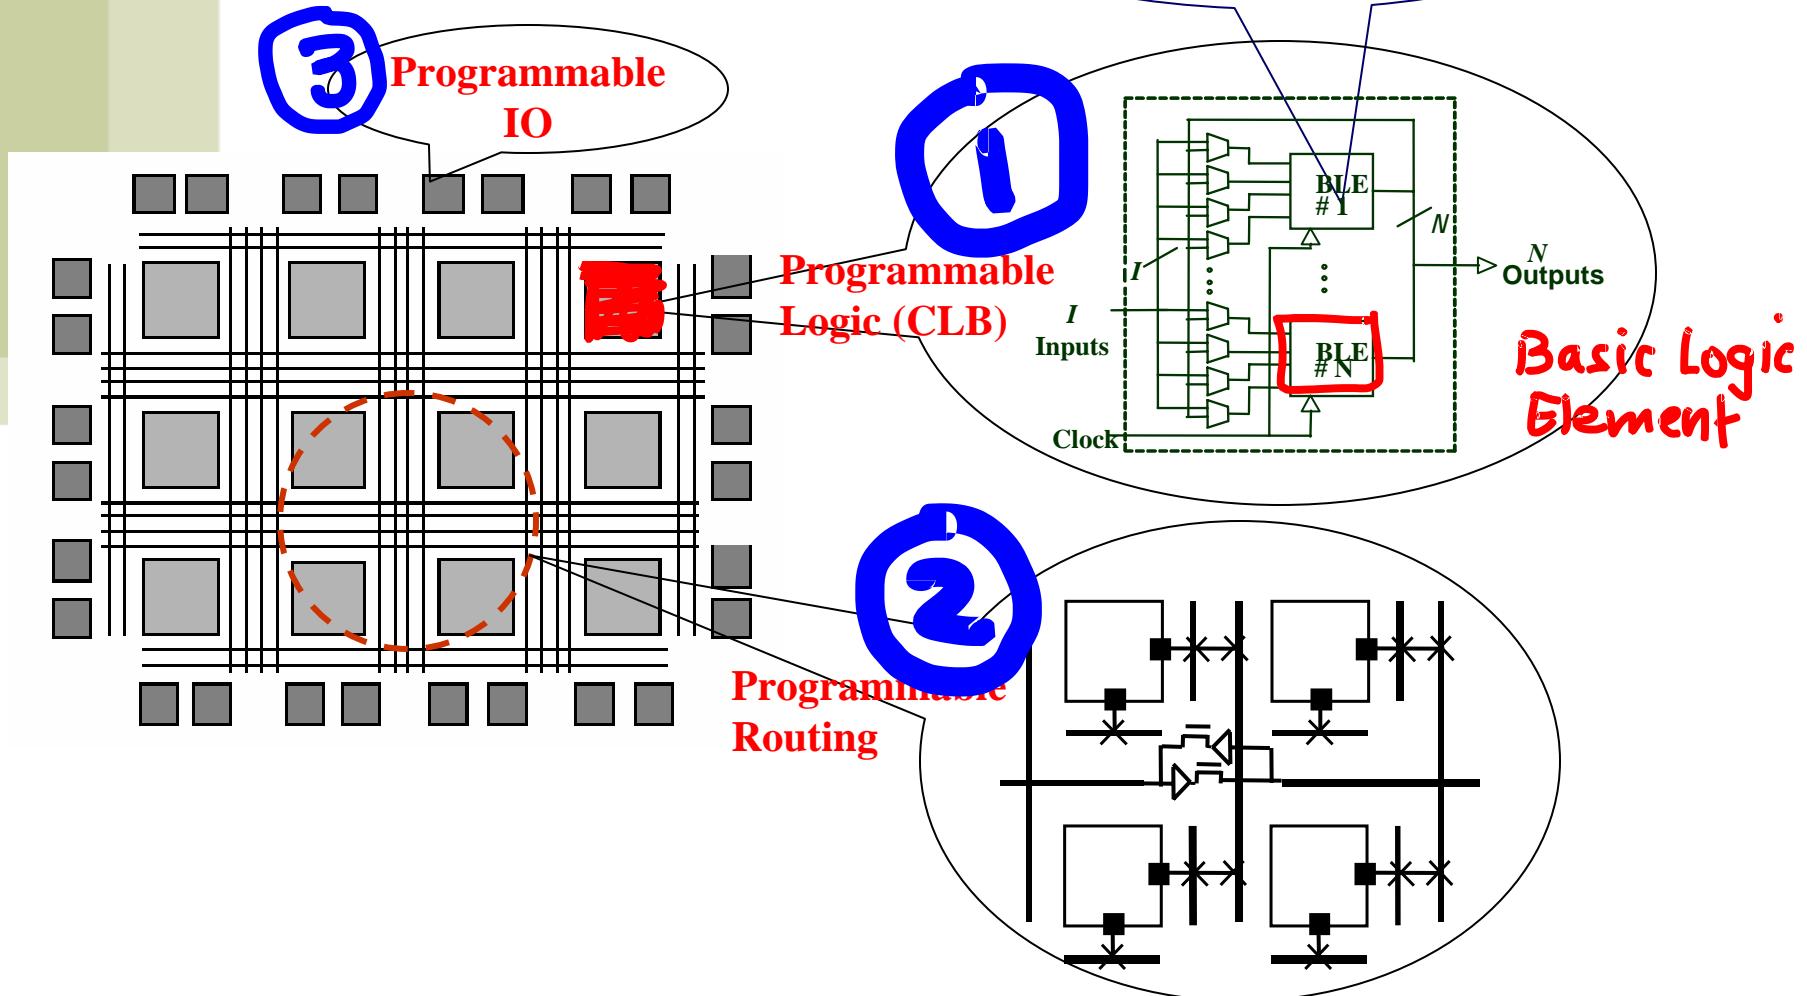
Field-Programmable Gate Arrays

- Based on Configurable Logic Blocks (CLB)



CLB } xilinx
BLE }
ALM Altera -

A Generic FPGA Architecture



(21.1)

The circuits in an FPGA for programming are based on MOSFETs (Metal Oxide Semiconductor Field Effect Transistor)
 aka transistors
 aka MOS devices
 aka FETs

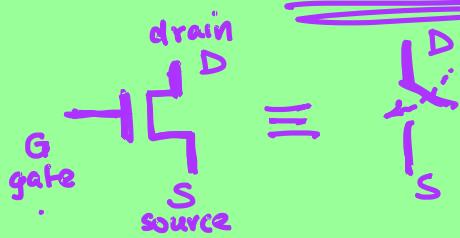
There are 2 kinds of MOSFETs :

PMOSFETs

NMOSFETs

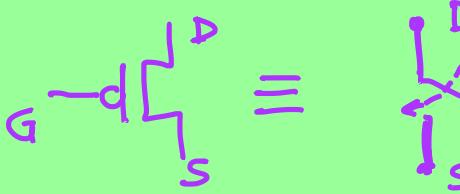
Both are switches (electrically)

(N)

 \equiv 

if $V_G = VDD$, the switch is closed
 (D & S shorted)
 else switch is open
 (D & S open ckt)

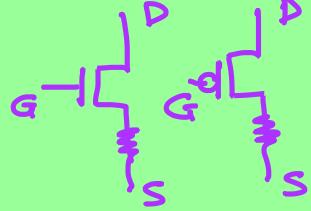
(P)

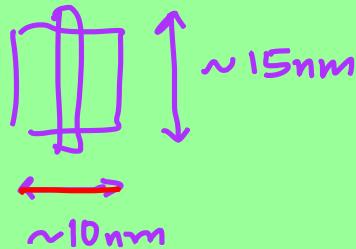
 \equiv 

if $V_G = GND$, the switch is closed
 else switch is open

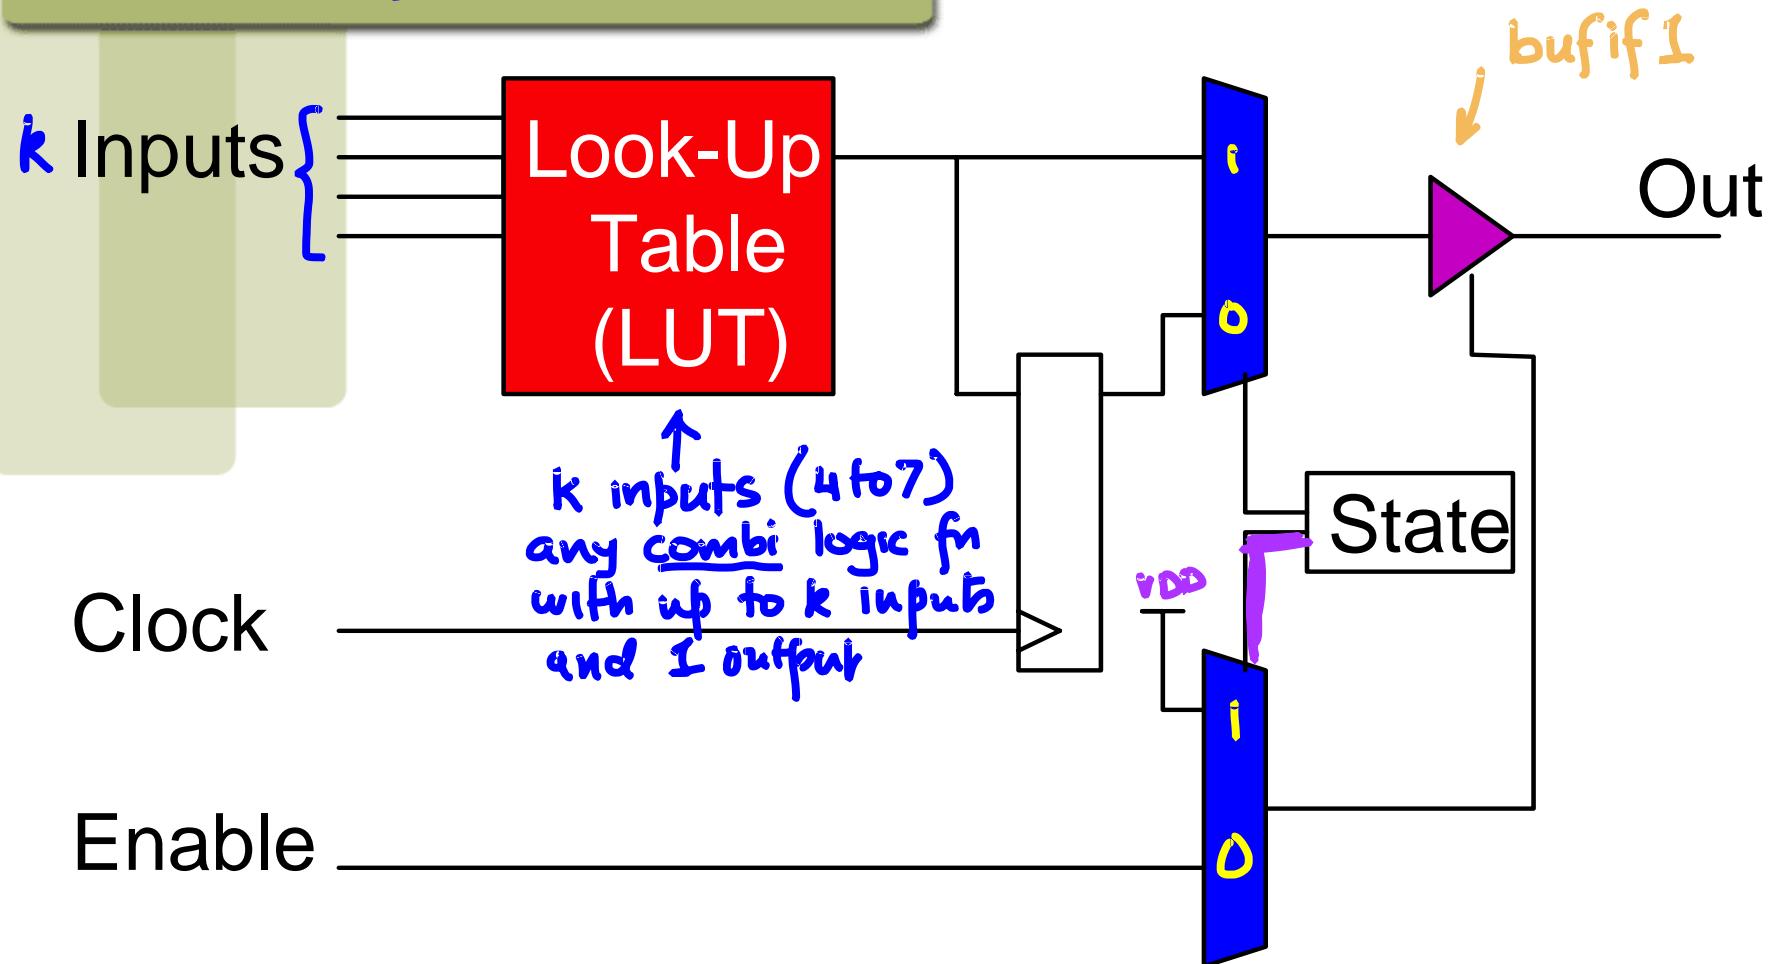
(21.2)

In reality

- PMOS & NMOS FETs have a series D-S resistance (not short circuited !)

- PMOS/NMOS device have on-resistance $\sim 5\text{ k}\Omega$ and off resistance of $> 1\text{ G}\Omega$
- They are tiny ("5nm technology"
has P or N devices slightly larger than 5nm in each dimension



What's in a ~~CLB~~? **BLE**



Xilinx CLB – a.k.a. “Slice”

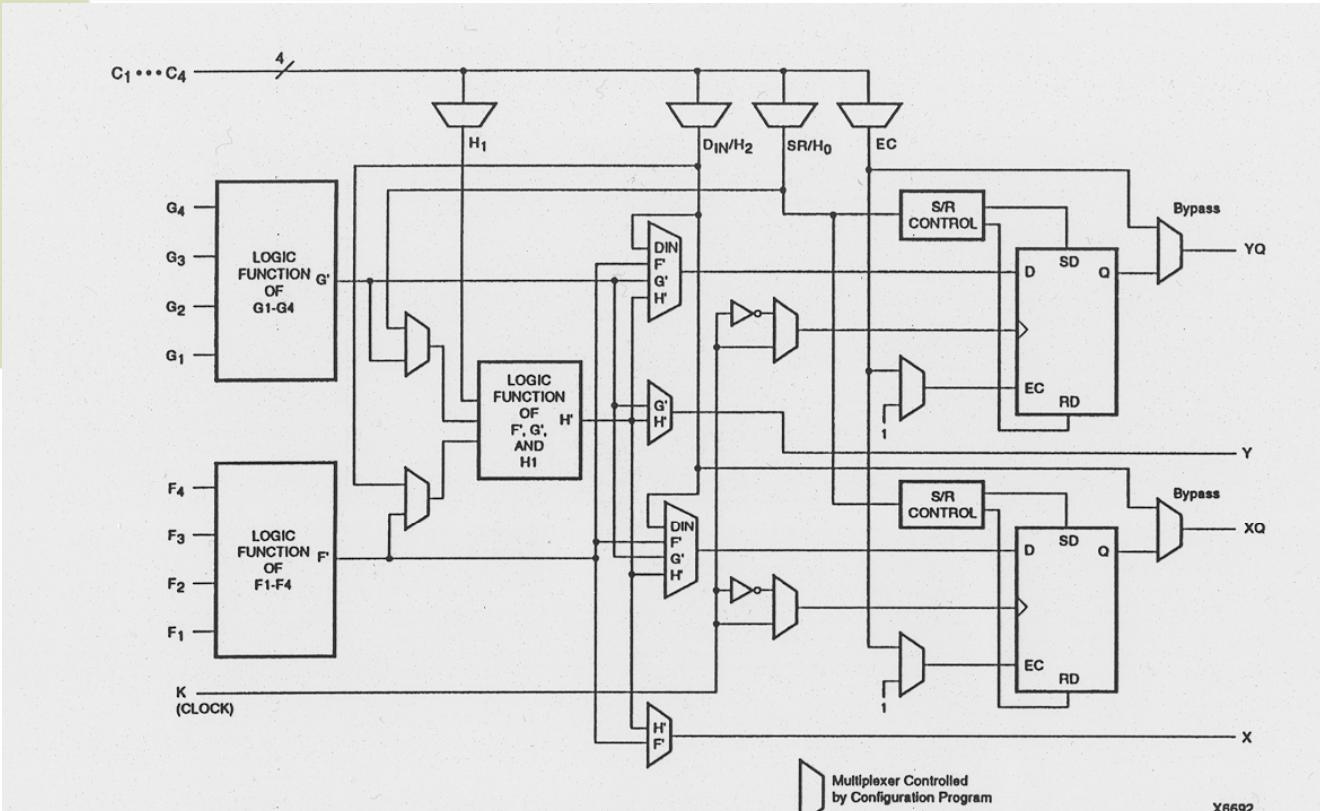
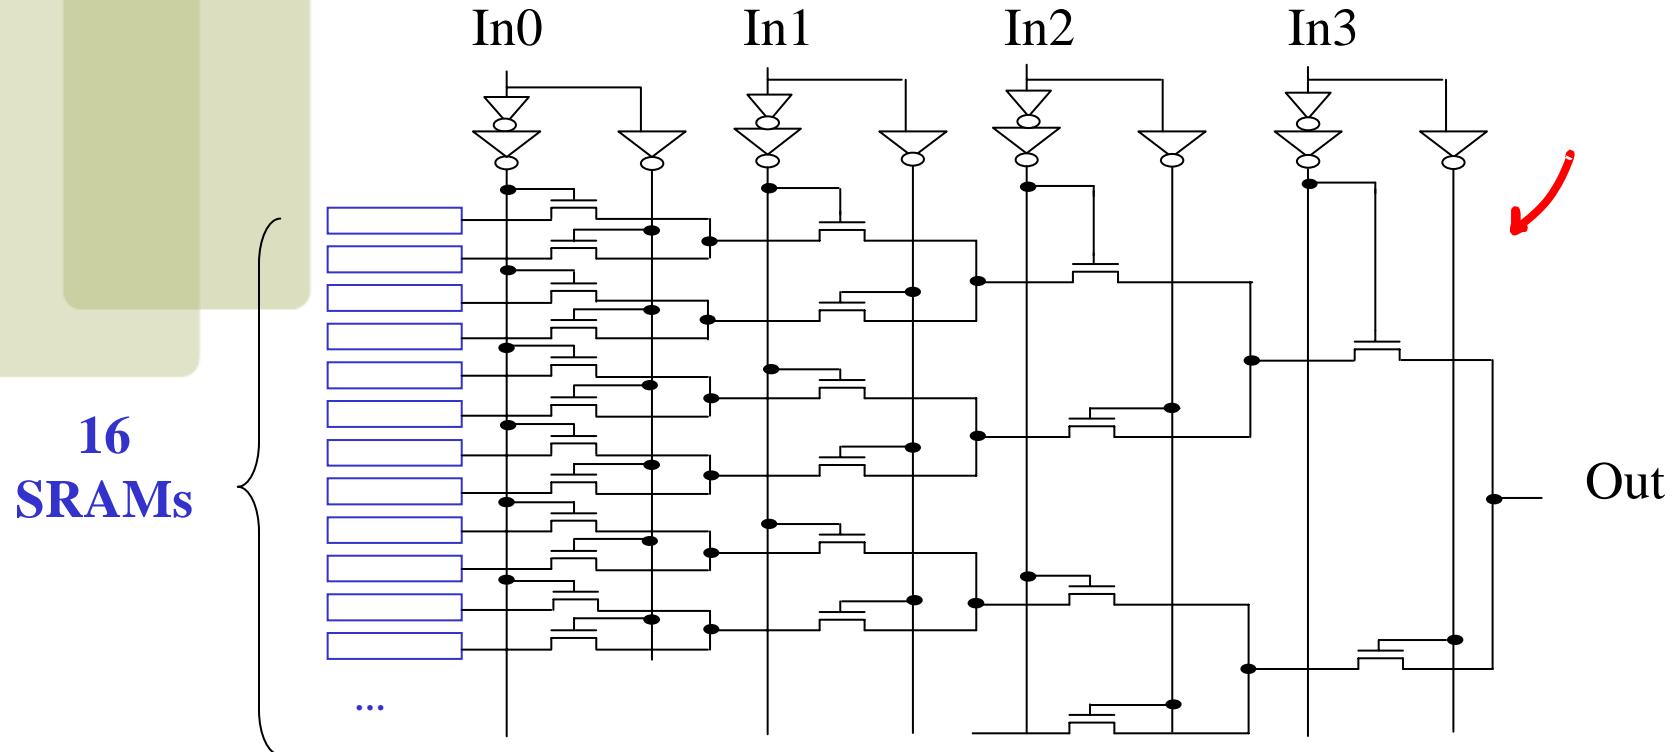


Figure 1: Simplified Block Diagram of XC4000-Series CLB (RAM and Carry Logic functions not shown)

Page 4-12, Xilinx XC400 Series Field Programmable Gate Arrays Product Specification

An Implementation of a 4-input Look-up Table (4-LUT)



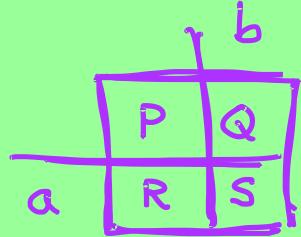
$$Out = f(in_0, in_1, in_2, in_3)$$

(24.1)

2-input LUT

- Implements any function of ≤ 2 inputs

•



If $P=Q=R=0, S=1$
 we get AND gate

0	0
0	1

- # of squares in K-map = $2^2 = 4$

[if n -inputs, then # of squares in Kmap is 2^n]

- # of possible functions of ≤ 2 inputs?
 $= 2^{\# \text{ of squares in Kmap}}$
 $= 2^2 = 2^4 = 16$

[if n -inputs then # of possible functions
 is $2^{\# \text{ of squares in Kmap}} = 2^{2^n}$]

24.2

Current LUTs use 4-6 inputs
functions they can implement

↳ 4: $2^{2^4} = 2^{16} = 64K$

5: $2^{2^5} = 2^{32} = 4G$

6: $2^{2^6} = 2^{64} = 64 \times \underbrace{2^{30}}_{1G} \times \underbrace{2^{30}}_{1G}$ WOW

K)

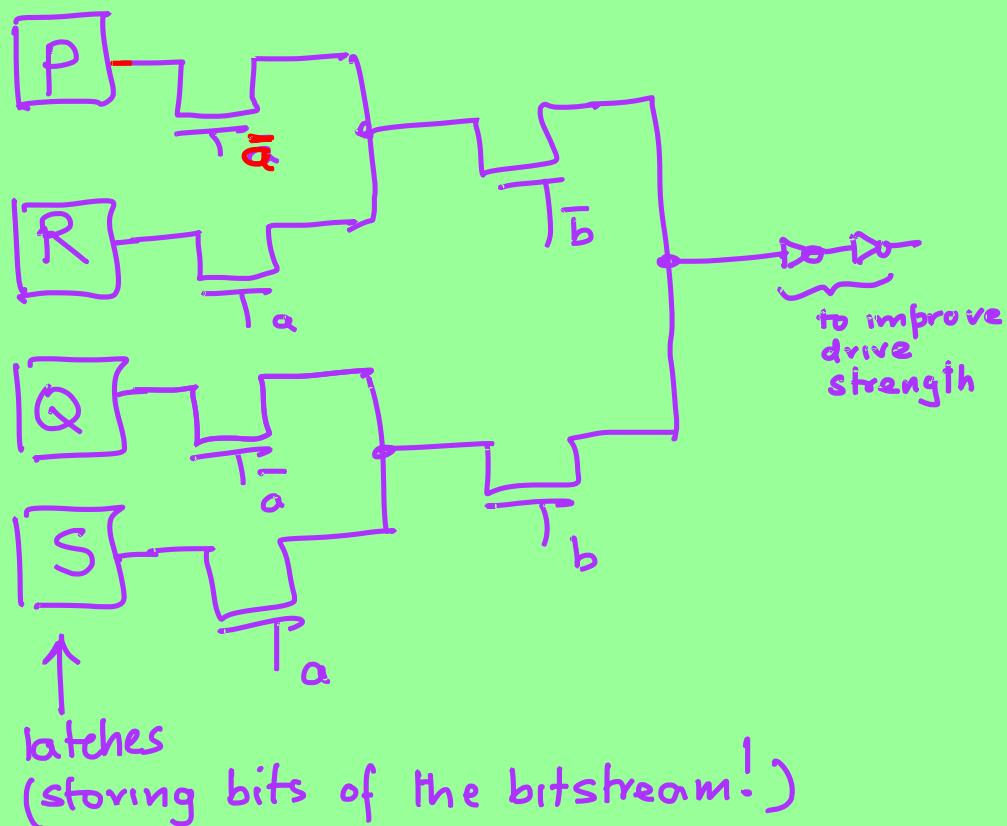
7: $2^{2^7} = 2^{128} = 256 G G G G$

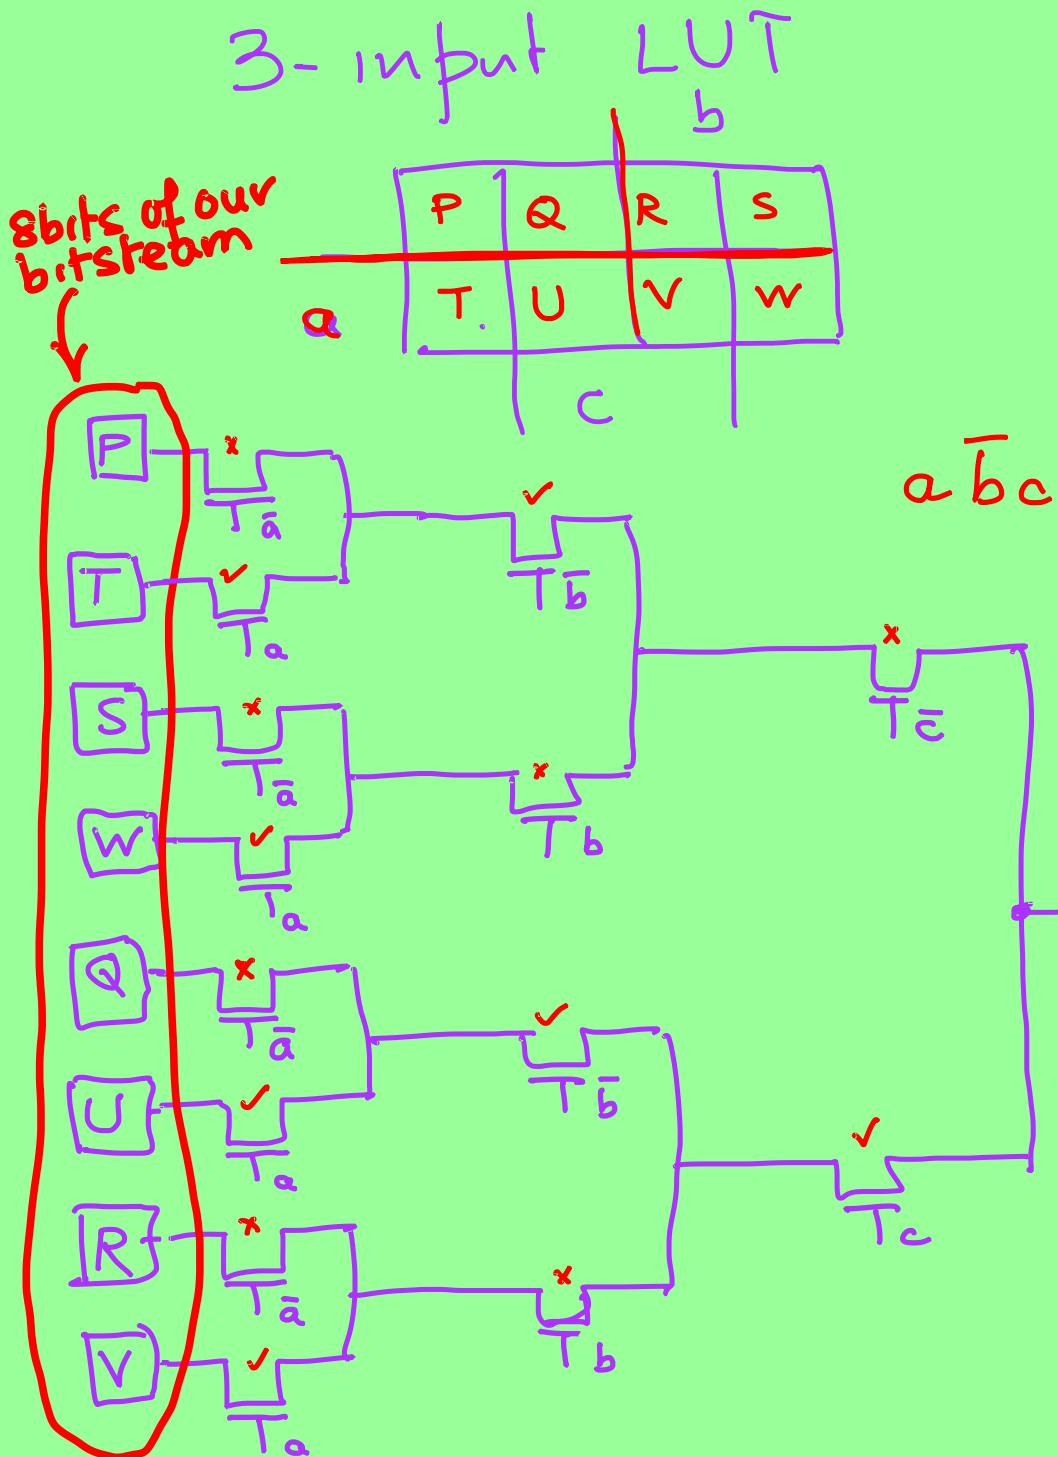
24.3

Lets create a 2-input
LUT:

	b
a	P Q
	R S

"if $\bar{a}\bar{b}$ then P
 elseif $a\bar{b}$ then R
 elseif $\bar{a}b$ then Q
 else if ab then S"





(24.4)

(24.5)

# inputs	# device	# functions
2	$4+2=6$	$2^2 = 16$
3	$8+4+2=14$	$2^3 = 256$
4	$16+14=30$	$2^4 = 64K$
5	$32+30=62$	$2^5 = 4G$
6	$64+62=126$	$2^6 = 16G\cdot G$

w/o counting
output \rightarrow

cost < benefit

Input-Output Blocks

- One IOB per FPGA pin
 - Allows pin to be used as input, output or bidirectional (tri-state)
- Inputs
 - Direct
 - Registered
 - Drive dedicated decoder logic for address recognition
- IOB may also include logic for boundary scan (JTAG)

Xilinx IOB

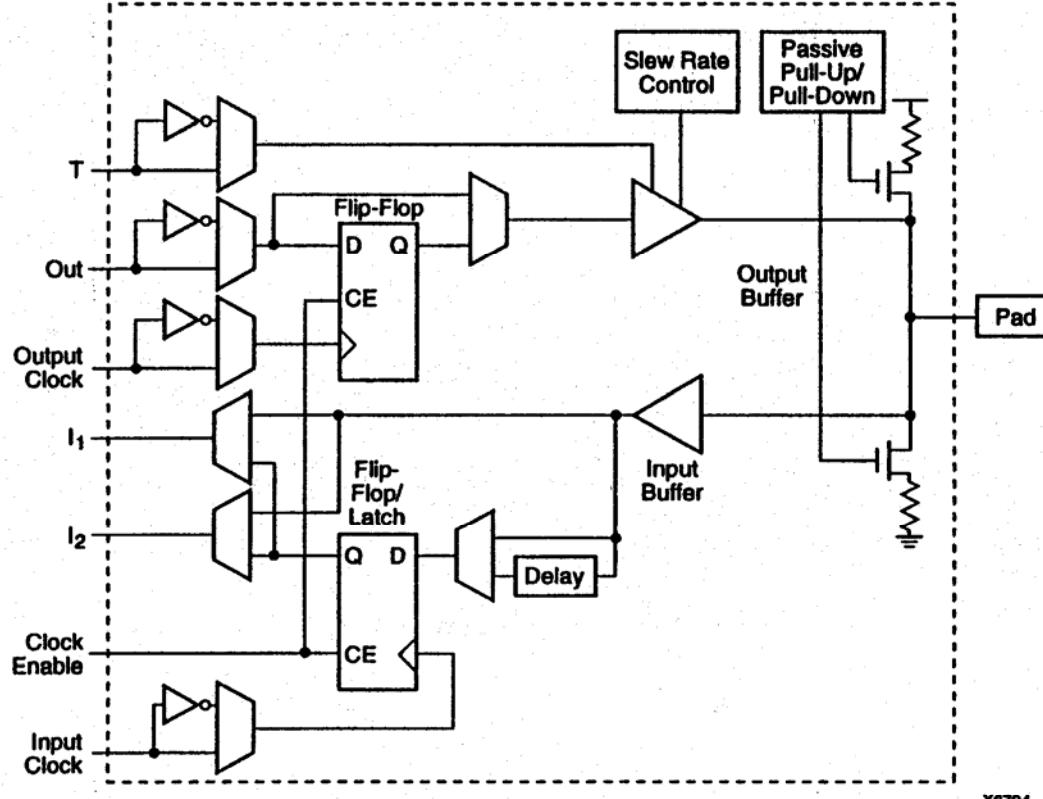


Figure 16: Simplified Block Diagram of XC4000E IOB

X6704

Figure 16, Xilinx XC400 Series Field Programmable Gate Arrays Product Specification

Interconnect

- 2-Dimensional mesh of wires, with switching elements at wire crossings to control routing
 - Bit patterns stored into the switch FFs determine routing
 - Switch connections programmed as part of configuring array
- To optimize for speed, many designs include multiple lengths of wire
 - Single-length (connect adjacent switches)
 - Double-length (connect to switches two hops away)
 - Long lines (run entire length/width of array)

Interconnect Diagram

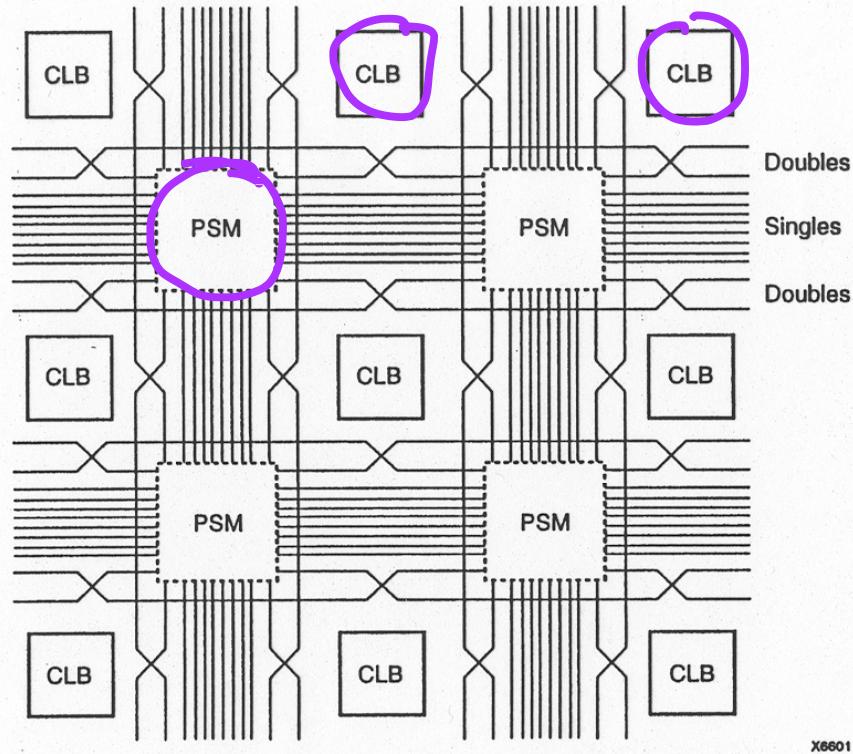
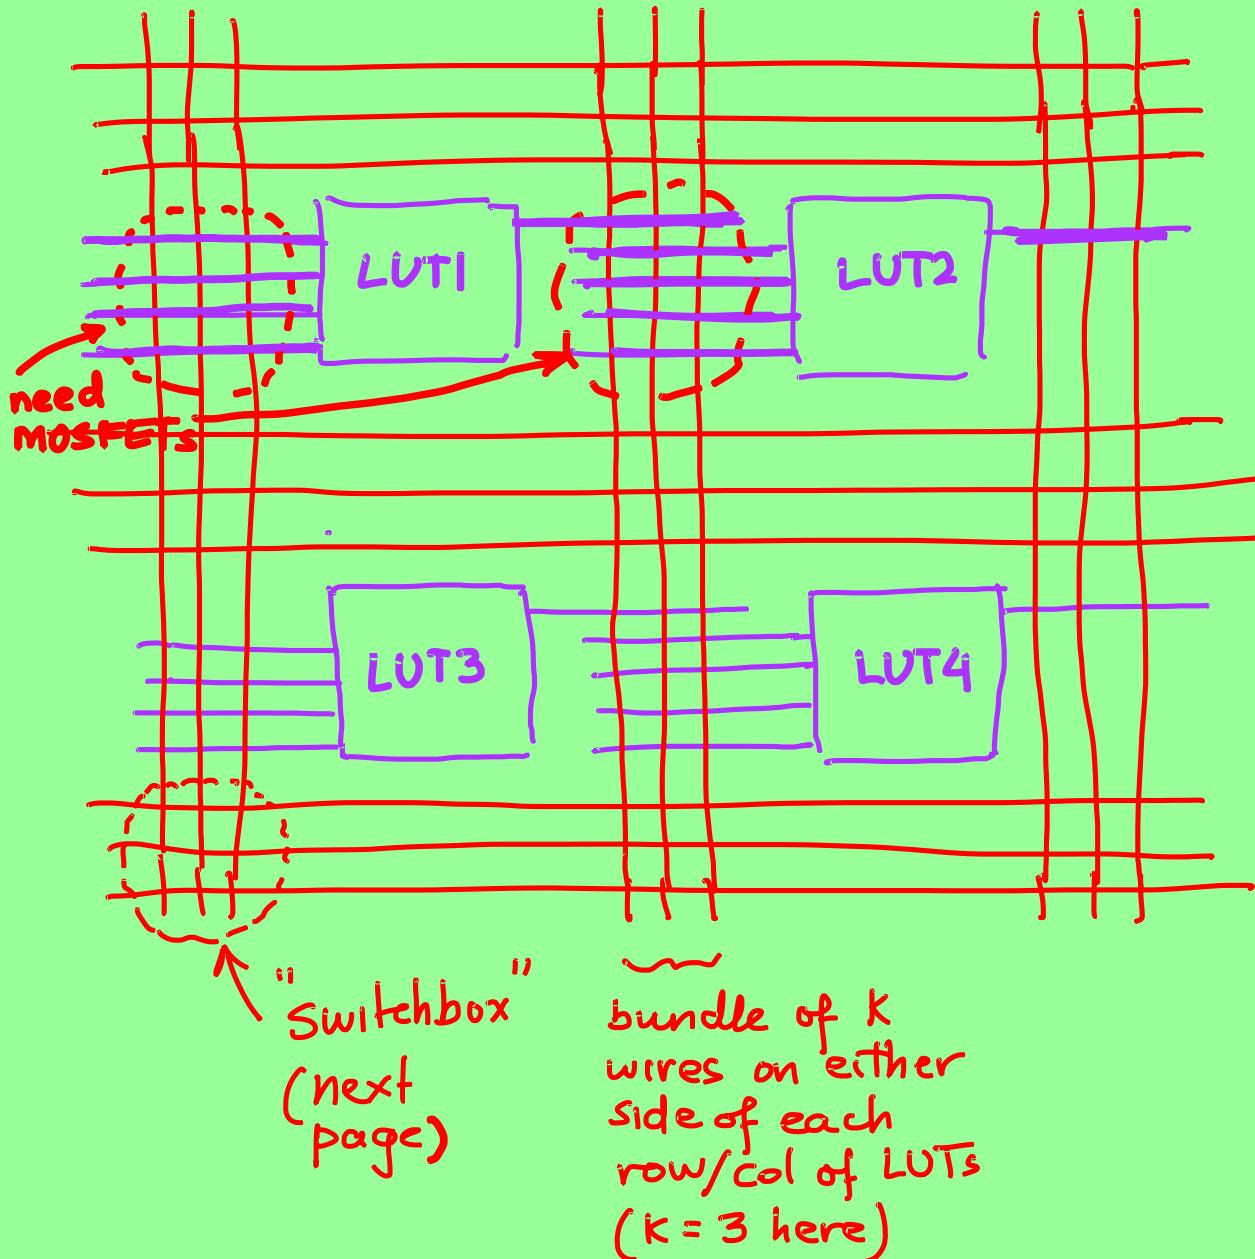


Figure 29: Single- and Double-Length Lines, with Programmable Switch Matrices (PSMs)

Figure 29, Xilinx XC400 Series Field Programmable Gate Arrays Product Specification

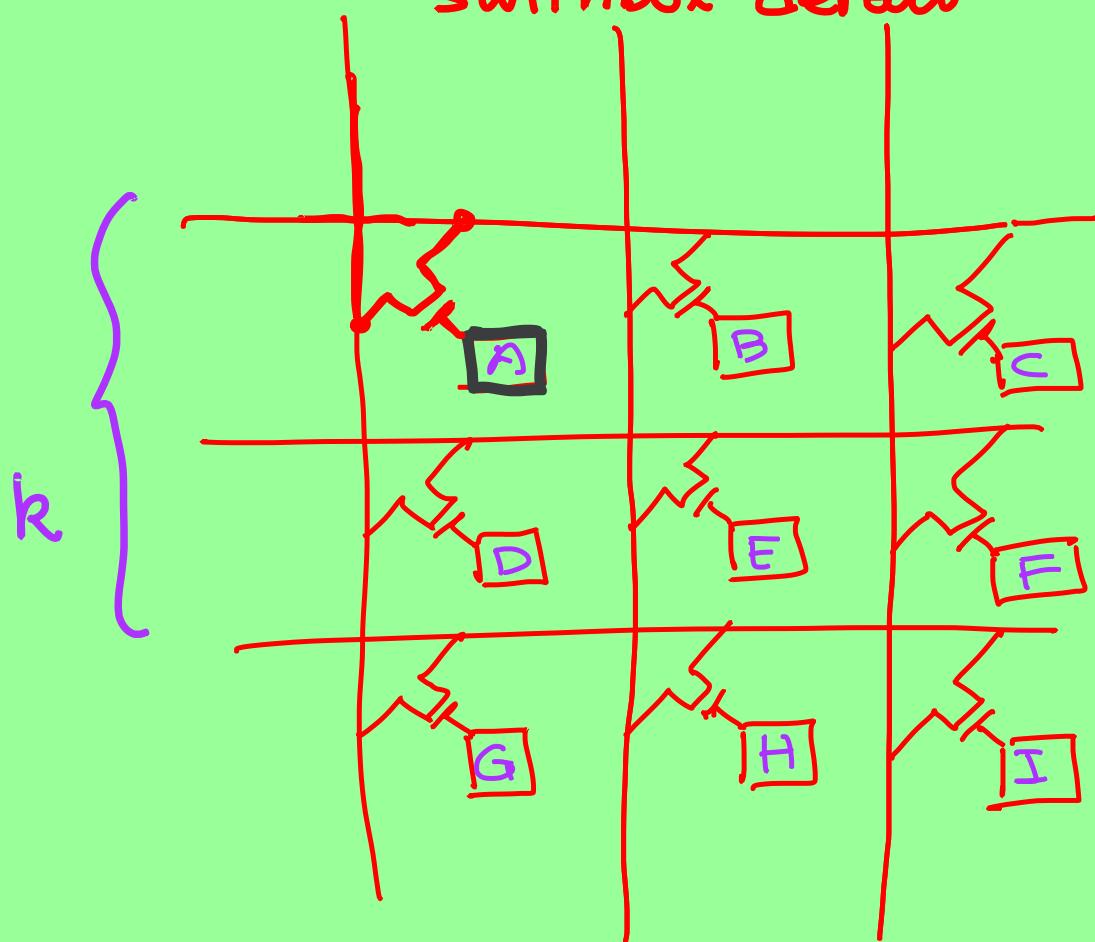
Wiring : Consider an FPGA with
4 LUTs (each 4-input)

(28.1)



(28.2)

switchbox detail



A, B, C, D, \dots, I are K^2
(9 in our example)
bits of our bitstream

Call 2L.L.1

(28.3)

Consider FPGA with 160000 5-input LUTs. Switchboxes size is 10×10 . The FPGA has 200 I/O pins, which can be set in one of 12 ways. How long is the bitstream?

$$I/O : [\log_2 12] \times 200 = 860b$$

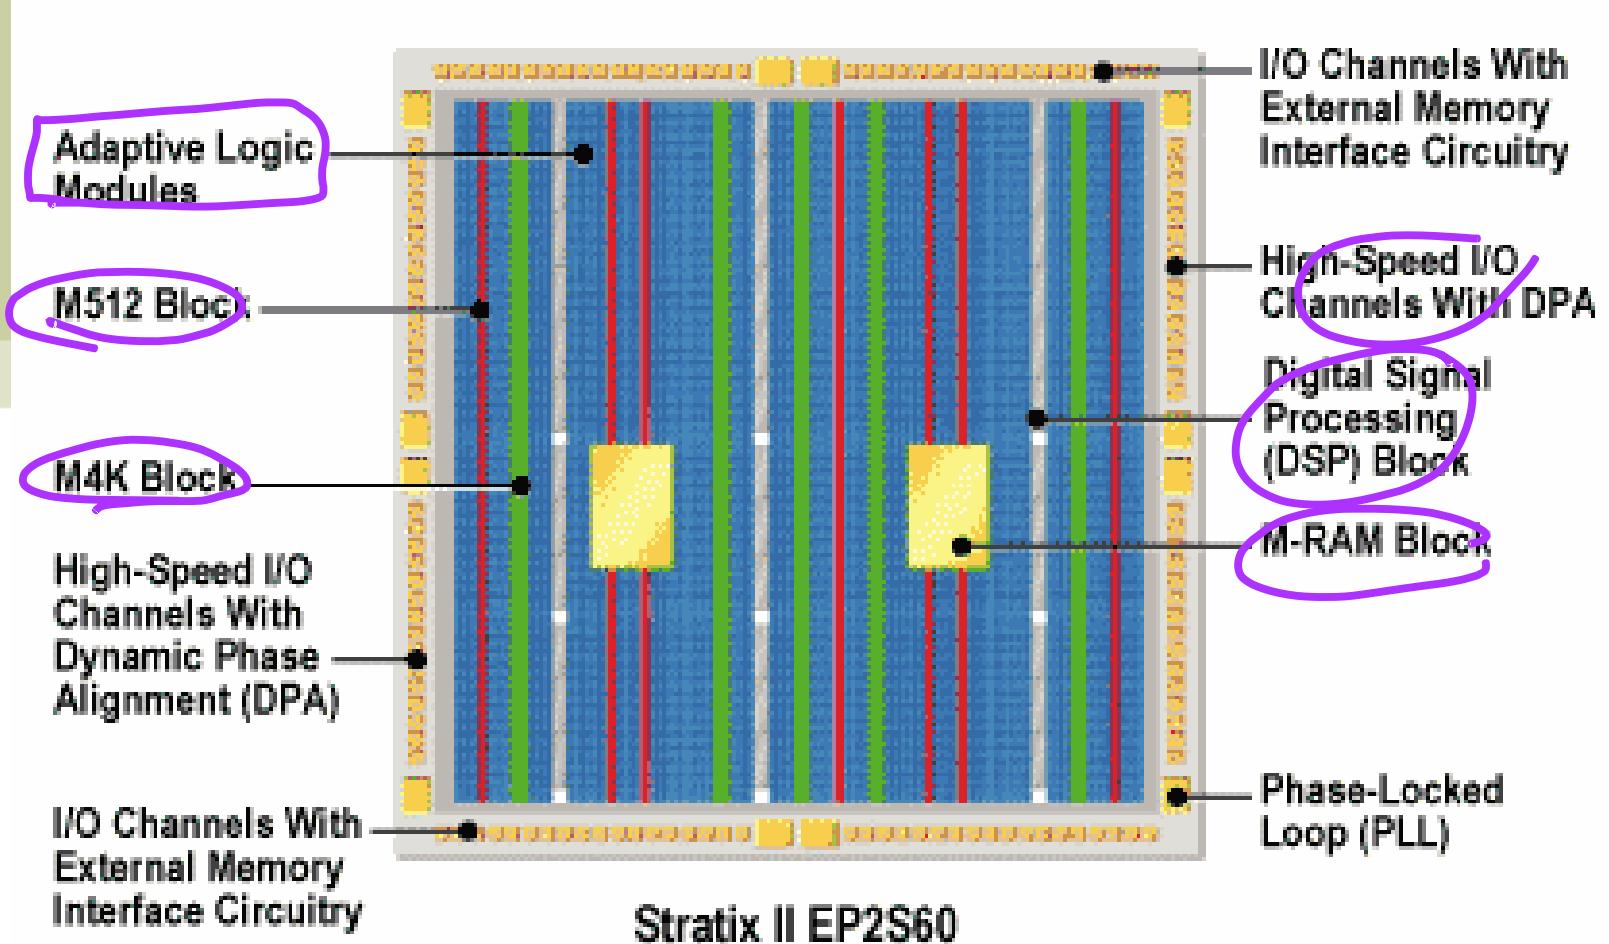
$$\text{LUTs} : (2^5 + 2) \times 160000 = 5.44 \times 10^6 b$$

Wires: The FPGA is 400×400 LUTs,

$$\text{Switchbox} = \underbrace{(401) \times (401)}_{\# \text{switchbox}} \times 100 = 16.08 \times 10^6$$

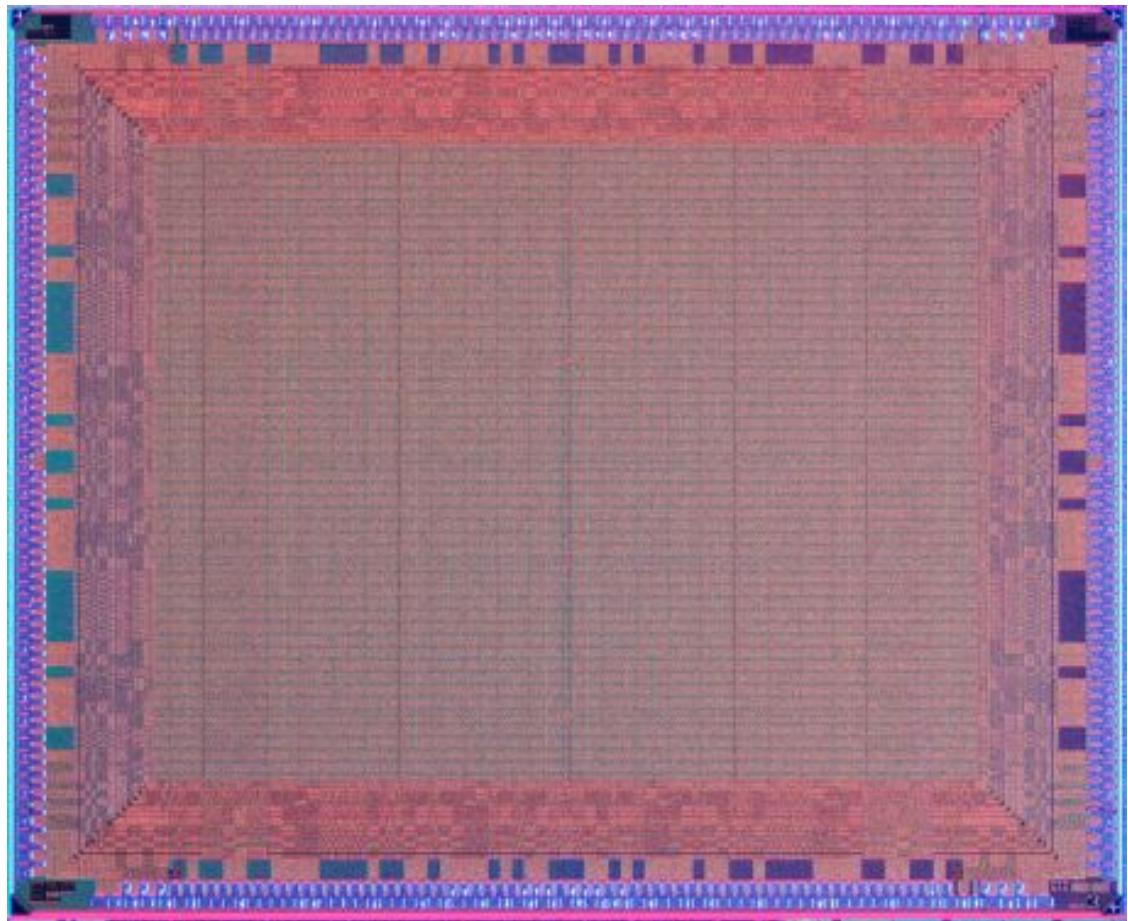
$$\text{Wires to LUTs} = \underbrace{(160000)}_{\# \text{LUT}} \times \underbrace{(5+1)}_{5 \text{ in } \& \text{ 1 out}} \times 10 = 9.6 \times 10^6$$

One Commercial FPGA, Altera Stratix II

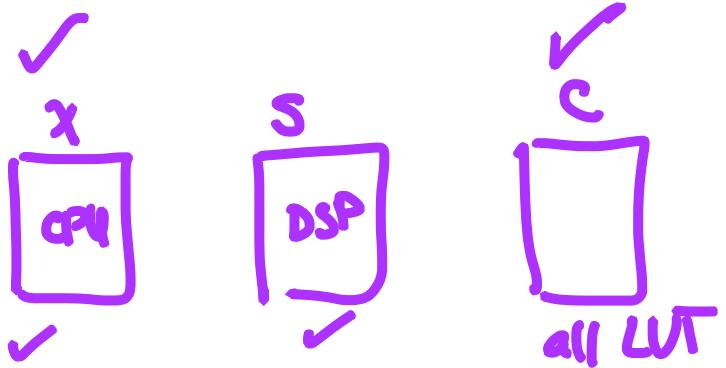


Chip Shot - Xilinx Spartan-3 die image

- Note the regularity...



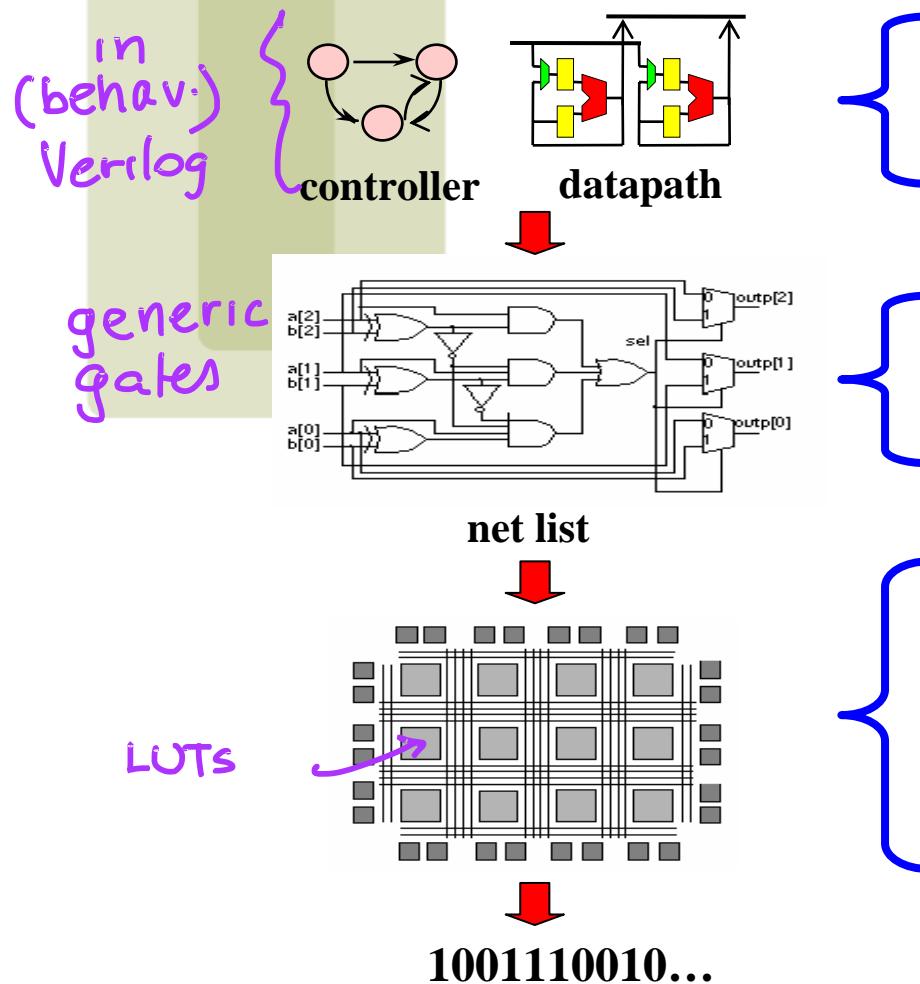
Design Variables



- The following issues are something that the company which designs the FPGA needs to worry about. The user of the FPGA is agnostic to these issues.
- # of inputs to LUT
 - Trade off number of CLBs required vs. size of CLB and routing area
- How is logic implemented
 - Switch based? Gate based?
 - SRAM configuration? Fuse burning configuration?
- Flip-flop in CLB?
- Additional Functionality
 - Carry chains, CPU's, block RAM files

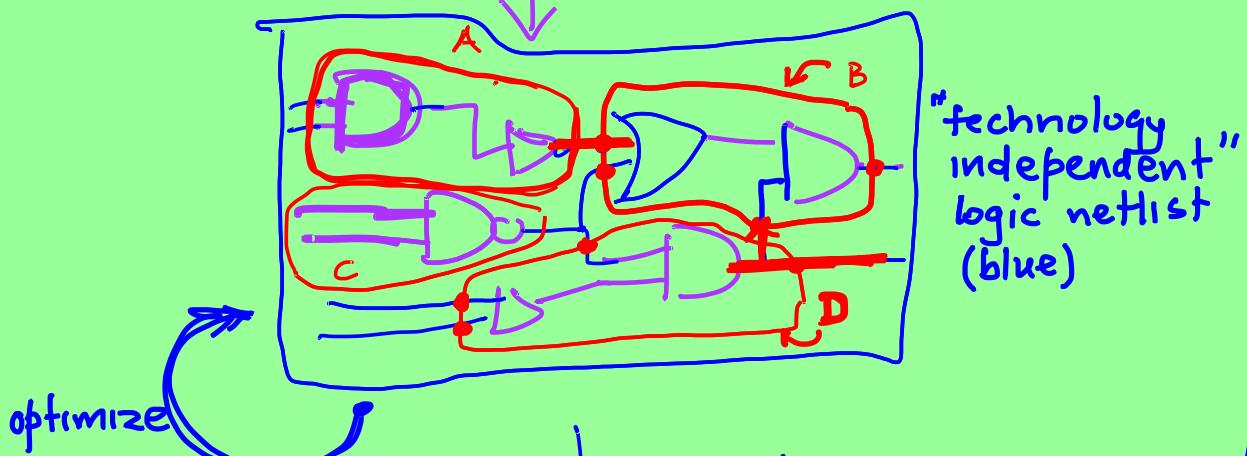
CAD flow

Design Flow for Programmable Logic



32.1

always@(a or b)
RTL
(Verilog)

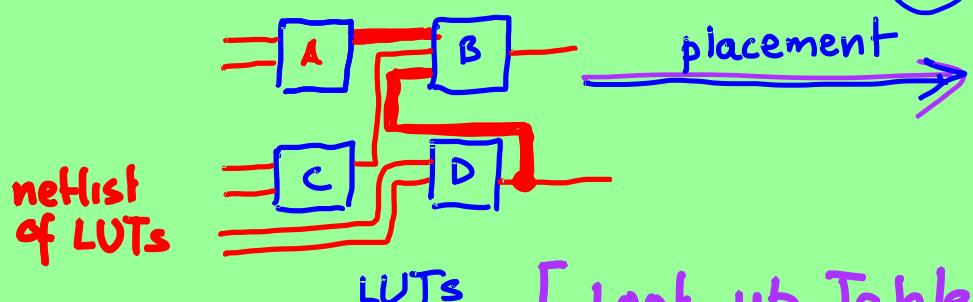
Logic
synthesis ①

"technology
independent"
logic netlist
(blue)

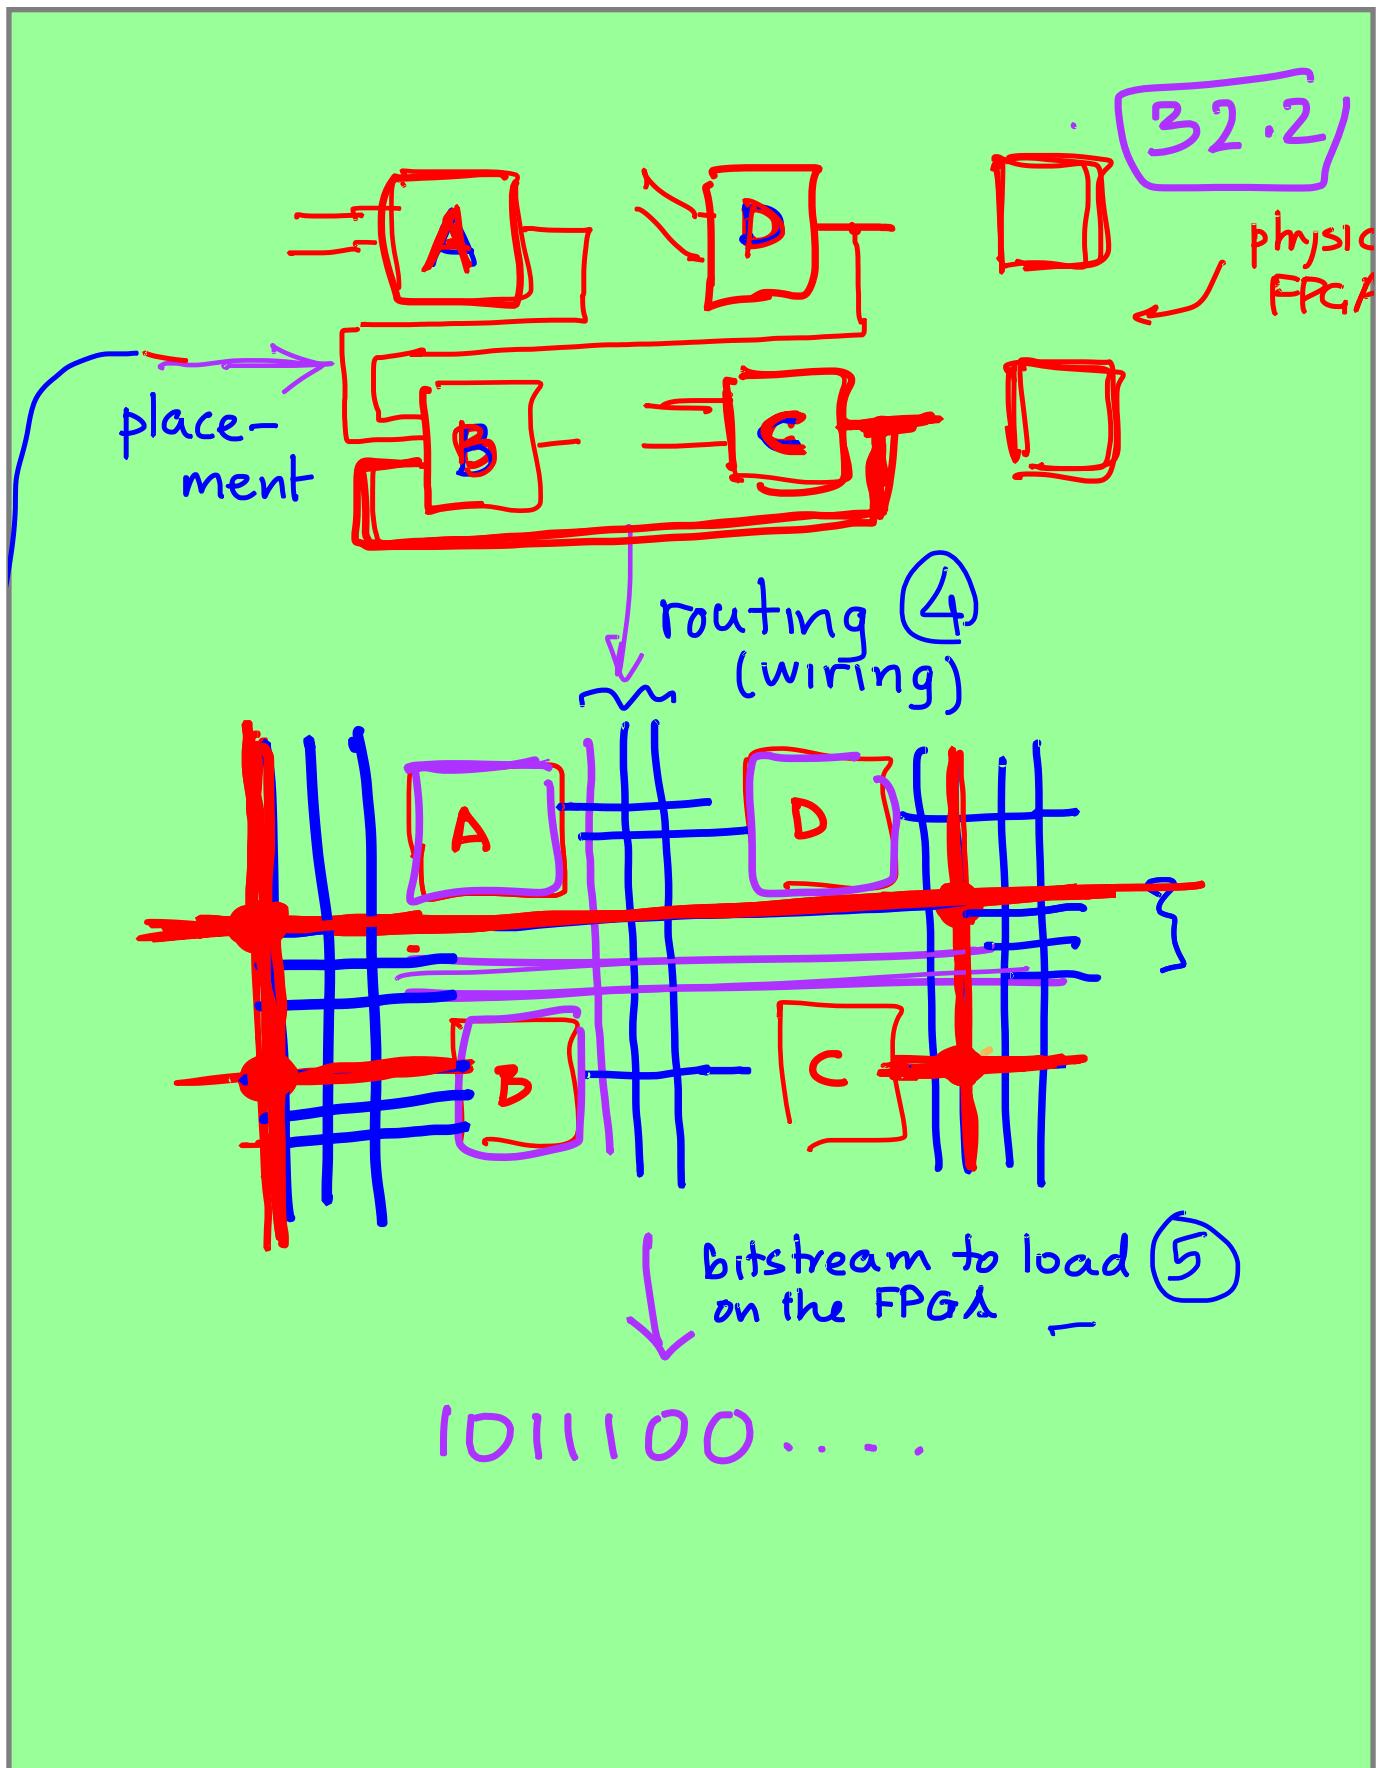
technology
mapping
(say $k = 3$)

②

③



[Look up Table]
any logic fn of $\leq k$ imp.



32.3

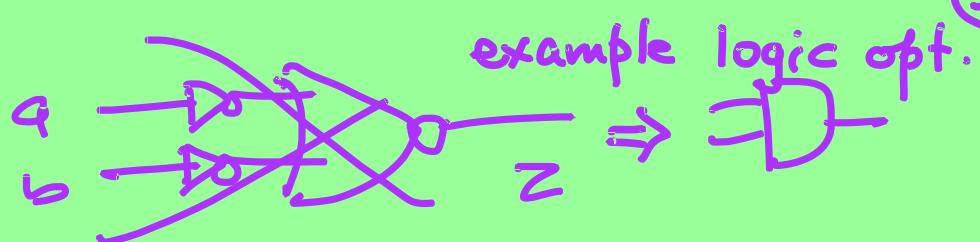
3 inp = $2^3 = 8$ boxes in kmap
 $2^8 = 2^{2^3}$ = #ways to fill box

4 inp = $2^4 = 16$...
; $= 2^{16} = 2^{2^4}$ = unique for

k inp = 2^k ...
 2^{2^k} unique for

6inp

 $2^{2^6} = 2^{64}$ functions



$$\overline{\overline{a} + \overline{b}} = z$$

$$z = ab$$

FPGAs – Pros (recap)

- Reasonably cheap at low volume
 - Good for low-volume parts, more expensive than IC for high-volume parts
 - Can migrate from SRAM based to fuse based when volume ramps up
- Short Design Cycle
- Reprogrammable (~1 sec programming time)
 - Can download bug fix into units you've already shipped
- Large capacity (100 million gates or so, though we won't use any that big)
 - FPGAs in the lab are “rated” at ~1M gates for 30K LE's
- More flexible than PLDs -- can have internal state
- More compact than MSI/SSI

FPGAs – Cons (recap)

- Lower capacity, speed and higher power consumption than building an ASIC
 - Sub-optimal mapping of logic into CLB's – often 60% utilization
 - Much lower clock frequency than max CLB max toggle rate – often 40%
 - Less dense layout and placement and slower operation due to programmability
 - Overhead of configurable interconnect and logic blocks
- PLDs may be faster than FPGA for designs they can handle
- Need sophisticated tools to map design to FPGA. But the FPGA vendor typically provides these tools (at a cost).

FPGA Design Flow

- Now that we know what the circuit structure inside an FPGA is, lets see how we go about programming an FPGA.
- In other words, we will briefly cover the steps that we undertake between
 - The conception of a design idea
 - The decision-making step of whether an FPGA will be the correct hardware platform for the design
 - The design flow we follow to obtain an FPGA based hardware realization of this design.

From Concept to Circuit

- Need to specify your design and then implement it as a functioning system
- Trade-offs between time/cost and efficiency
 - Performance of final system
 - Amount of silicon area required (manufacturing cost)
 - Time to manufacture
 - Power consumption
- Need to think about a number of factors to make decision
 - Sales volume
 - Profit margin and how performance affects it
 - Time-to-market concerns, particularly if trying to be the first product in a new area

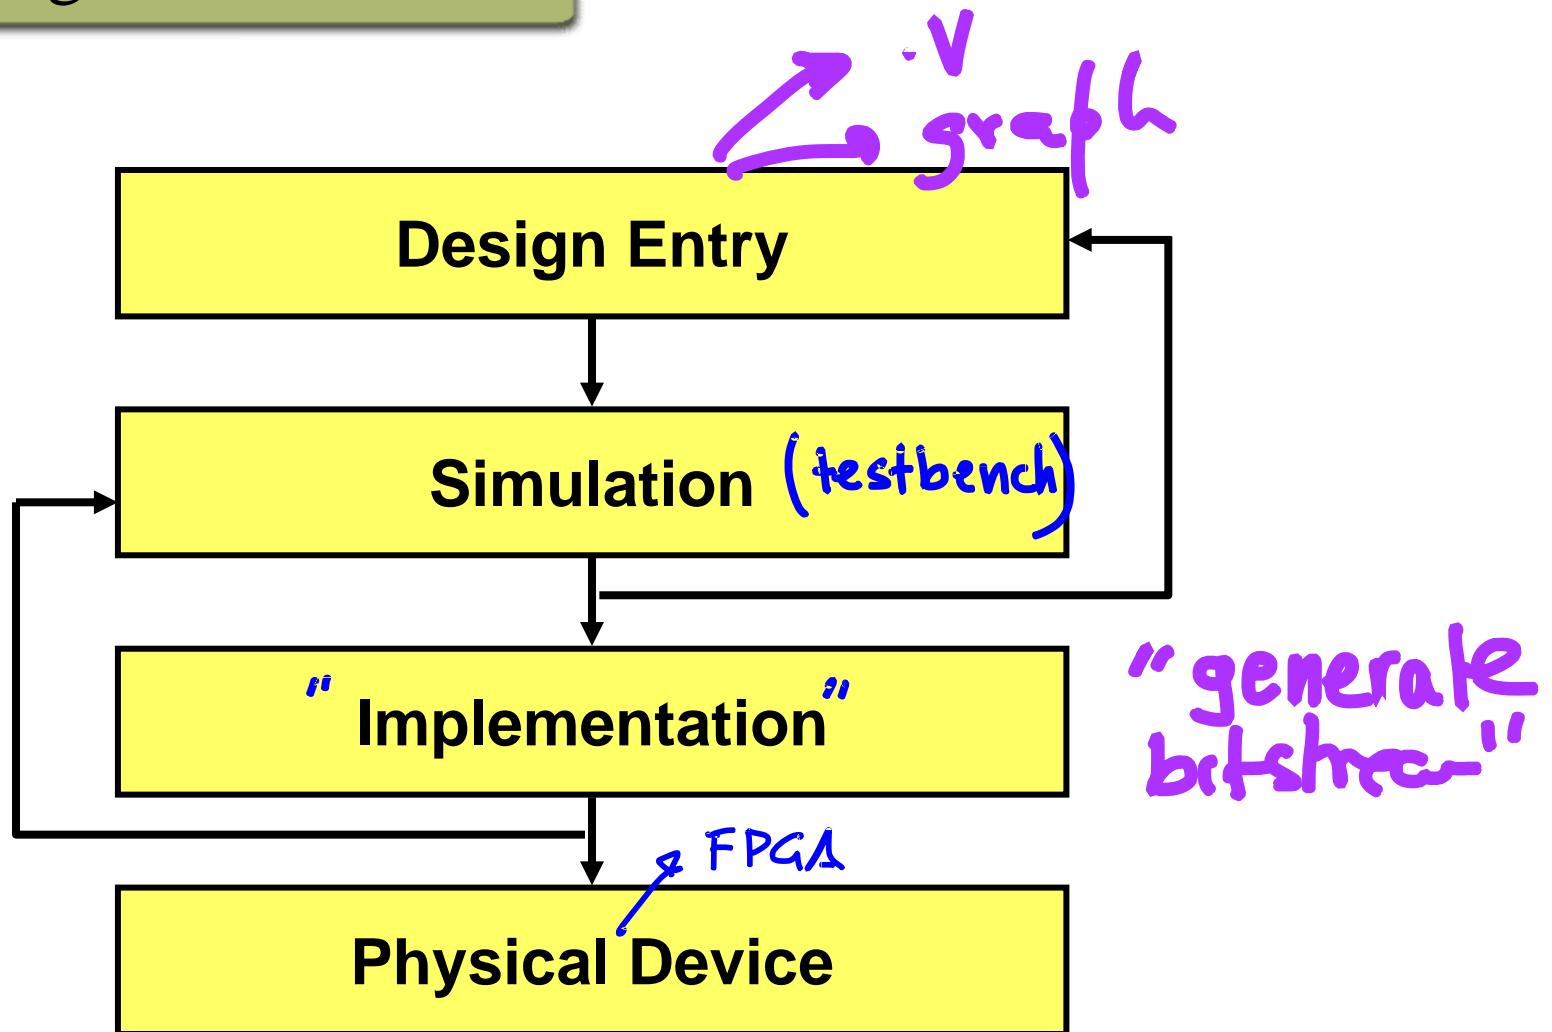
High-Level Design

- Problem 1: modern designs are just too complex to keep in your head at one time
 - Custom chips approaching 100M transistors
 - FPGA designs approaching 1M gates
- Problem 2: Even if you could design complex systems by hand, it would take too long
 - 100M transistors at 10s/transistor = 133.5 person-years
 - Transistor counts and system speed increasing at 50% or so/year
 - Design time is critical
- # transistors per chip increasing at 50%/year
- # transistors per engineer-day increasing at 10%/year
- Need techniques that reduce the amount of human effort required to design systems
 - Let humans work at higher levels, rely on software to map to low-level designs

Increasing Design Abstraction

- Old way: specify/layout each device by hand
 - Early chips were laid out by cutting patterns in rubylith with knives
- Current State of the Art: Combination of synthesis and hand design
 - Specify entire system in HDL (Verilog or VHDL), simulate, and test
 - Use synthesis tools to convert non-performance-critical parts of the design to transistors/gates
 - Human designs critical components by hand for performance
- Where Things are Going: System-on-a-Chip Design
 - Specify design out of high-level components (cores)
 - Integrate sensors, transmitters, actuators, computers on a chip
 - Rely very heavily on tools to map design to software and hardware.
 - ***XUP (the board we will use in the lab) is an SoC design vehicle***

(FPGA) Design Flow



Design Entry

Two main methods:

- Text entry (VHDL/Verilog):
 - Compact format, no special tools required
 - Good for high-level designs and control logic
- Schematic Capture: Draw pictorial representation of circuit, tool converts into design (typically HDL description)
 - Traditionally used for low-level (transistor) designs, regular structures
 - Commonly used today in conjunction with text entry to provide visual viewing of overall structure of a design

Simulation

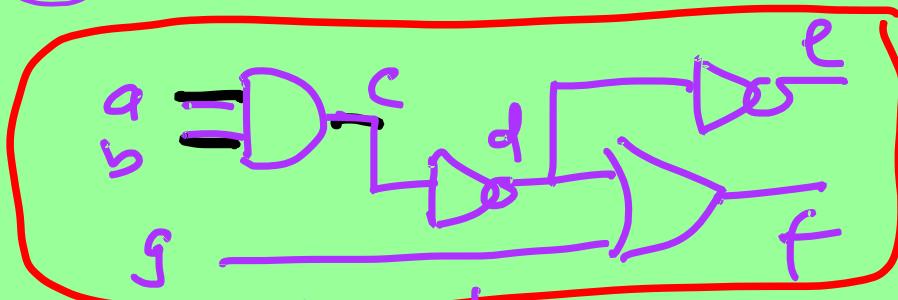
- Two types of HDL simulators
 - Interpreted: runs slower but more versatile and no compilation time
 - Compiled: runs faster but require compilation time and often not as versatile partly due to needs to compile all library components used.
- Both typically use Discrete-Event techniques
 - Divide time into discrete steps
 - User can select time step to trade accuracy vs. run-time
 - Keep lists of events that have to be resolved at each time step.
 - At each time step, resolve all events for the time step and schedule events for later time steps
- Output:
 - Text from output/print statements in your design
 - Errors from assert statements
 - Waveform traces
- Like any testing, the key is having good tests. The designer creates these!

event driven timing simulator

SIMULATION

41.1

A Interpreted
event driven



$\begin{array}{r} 1111 \\ + 10101100 \\ \hline 11010111 \\ - 10000100 \\ \hline \end{array}$

- read gate
- perform event driven simulation (see verilog notes) with or without timing

- N gates
- $T = \frac{1}{f}$ = CPU clock period that runs the event driven simulator
- α : activity factor
 - 20-25% (logic)
 - 90% arithmetic

$$\text{time to simulate} = \underbrace{100 \cdot T}_{\text{cycles per event}} \cdot \underbrace{\alpha}_{\text{\#events}} \cdot N \quad (\text{sec})$$

4T2

• Compiled code 32b words

c = and(a, b)
d = complement(c)
e = complement(d)
bitwise operator of ISA f = or(d, g)

↑ same ckt as on page 41.1

$$\text{time to simulate} = \frac{N}{32} \cdot T$$

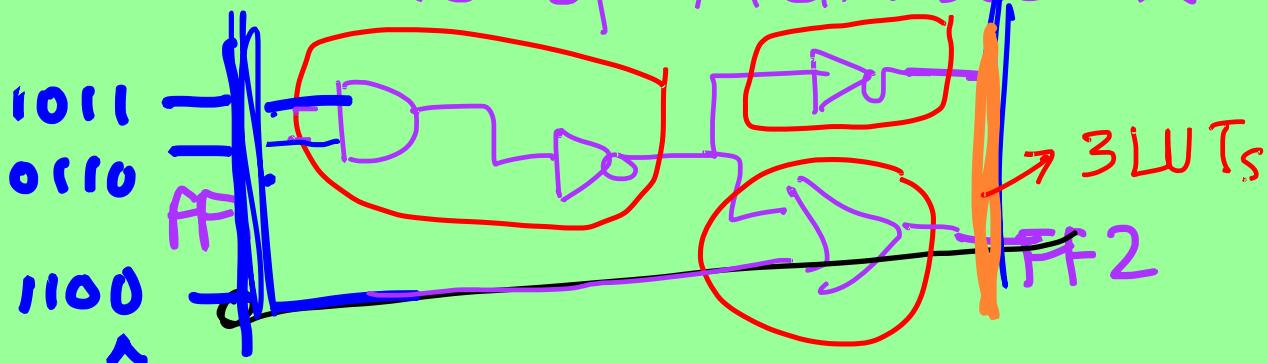
41.3

emulation (hardware)

(C)

10x20

take the netlist, map to FPGA & emulate on a "fridge-sized" rack of FPGA boards!



D = Logic depth (in LUTs) of the emulated circuit

τ = FPGA LUT delay

$$\text{time to simulate} = D \cdot \tau$$

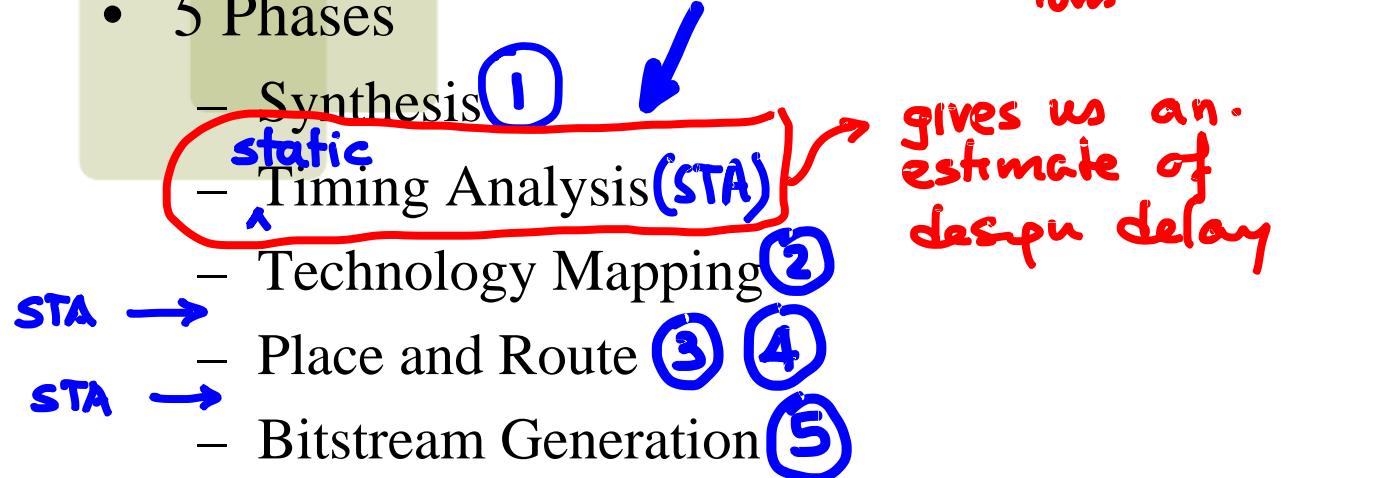
celaro

Rx

Implementation of an FPGA Design

Going from simulated Verilog design to circuits

- 5 Phases



$$T = 10\text{ns}$$
$$f = \frac{1}{T} = 100\text{MHz}$$

(Sometimes do additional timing analysis after place and route just to make sure that the timing is good)

Synthesis

Verilog \leftrightarrow

Transforms program-like VHDL into hardware design (netlist)

- Inputs
 - HDL description
 - Timing constraints (When outputs need to be ready, when inputs will be ready, data to estimate wire delay)
 - Technology to map to (list of available blocks and their size/timing information)
 - Information about design priorities (area vs. speed)
 - effort level

For big designs, will typically break into modules and synthesize each module separately

- 10K gates/module was reasonable size 5 years ago, tools can 50-100K gates now

Timing Analysis

Static timing analysis is the most commonly-used approach

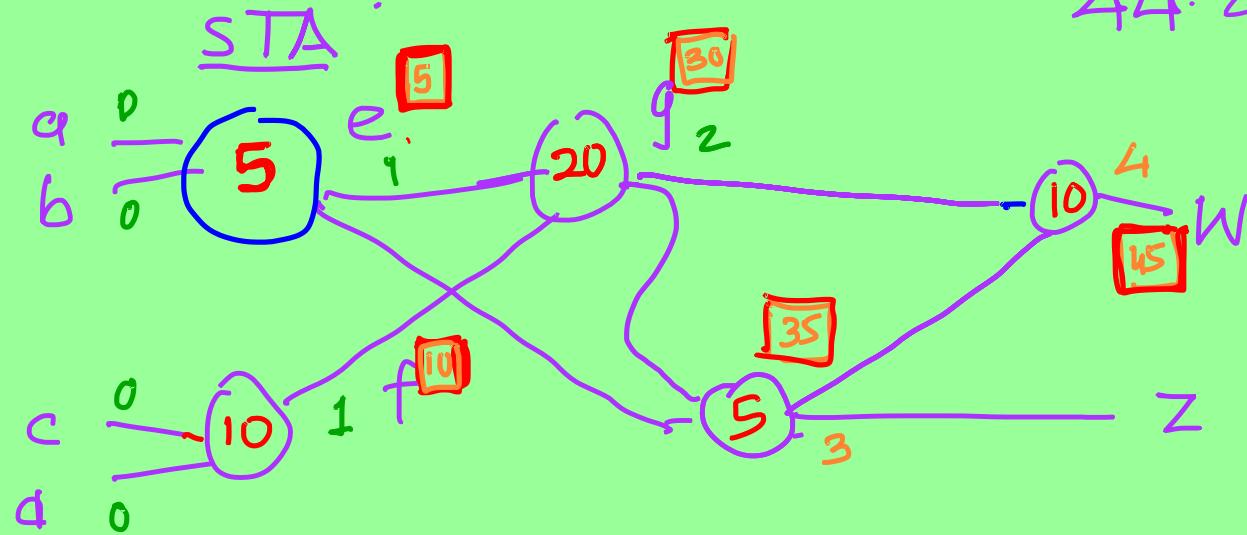
- Calculate delay from each input to each output of all devices
- Add up delays along each path through circuit to get critical path
- Works as long as no cycles in circuit
 - Tools let you break cycles at registers to handle feedback
- Also, ignores *false paths* in the design.
conservative O(n)
- Trade off some accuracy for run time
 - Simulation tools like SPICE will give more accurate numbers, but take much longer to run
- If the netlist passes timing analysis tests, we proceed further

44.1

Timing analysis

- { In: logic ckt (w/ N gate)
 delays of gates
Out: Max delay of ckt
- static timing analysis (STA)
 - $O(N)$
 - fast
 - conservative (overestimate the true delay)
- exact timing analysis
 - $O(2^N)$
 - exact
 - not viable in practice

>



① levelize gates (orange)

$$\cdot \text{level}(g) = 1 + \max_{\substack{f_i \in \text{fanin} \\ \text{of } g}} (\text{level}(f_i))$$

Some \rightarrow
gate

$$\cdot \text{level}_{(\text{primary})} = \emptyset$$

② Arrival time (AT) of a gate g (in level order)

$$\rightarrow \text{AT}(g) = \left[\max_{\substack{f_i \in \text{fanin} \\ \text{of } g}} \text{AT}(f_i) \right] + \Delta(g)$$

$$\text{AT}(e) = 5, \text{AT}(f) = 10$$

$$\text{AT}(g) = \max(5, 10) + 20 = 30$$

$$\text{AT}(z) = \max(5, 20) + 5 = 35$$

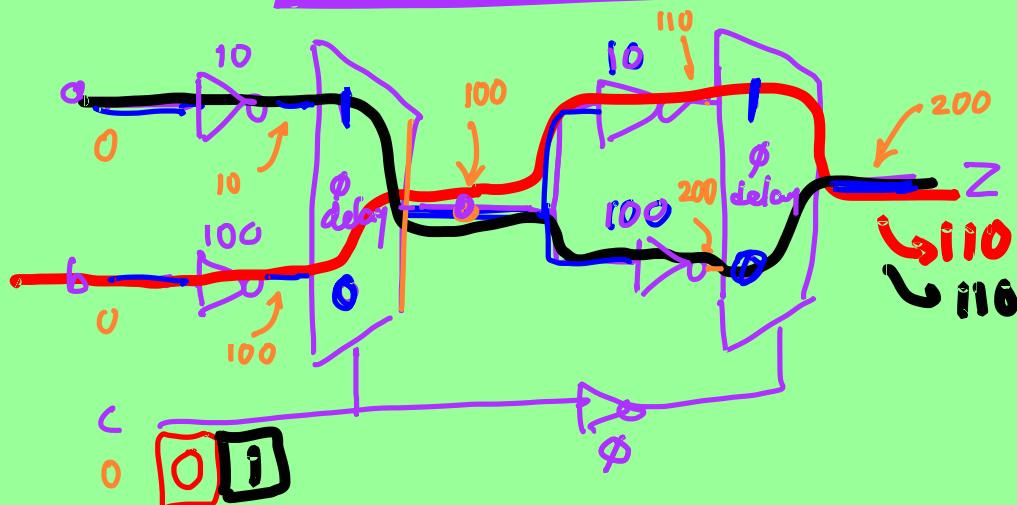
$$\text{AT}(w) = \max(30, 25) + 10 = 45$$

44.3

STA ^{run} time

$O(N)$, $N = \# \text{ gates}$

STA conservative pessimism



$$AT(z) = \underline{200}$$

[if real delay is Δ
then STA result will be $\geq \Delta$]

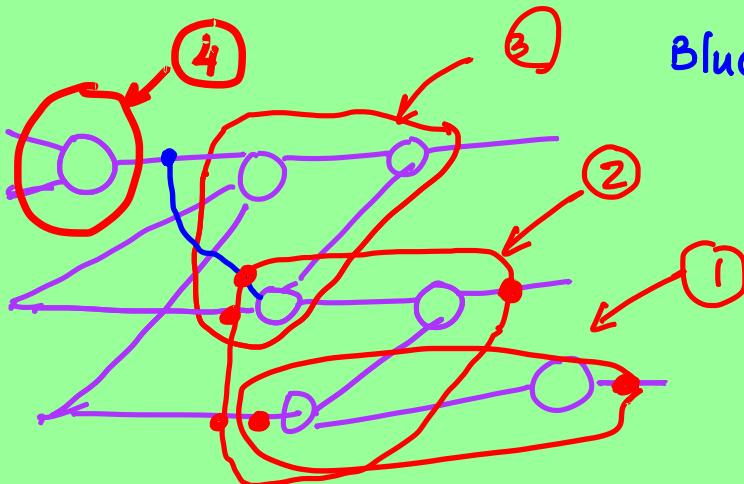
110
200

conservative

Technology Mapping

- Technology mapping converts a given Boolean circuit (a netlist) into a functionally equivalent network comprised only of LUTs or PLAs
 - Basically, can divide logic into n-input functions, map each onto a CLB.
- Technology mapping is a crucial optimization step in the programmable logic design flow
- Direct impact on
 - Delay (number of levels of logic)
 - area/power (number of LUTs or PLAs)
 - Interconnects (number of edges)
- Harder problem: Placing blocks to minimize communication, particularly when using carry chains

45

Tech Mapping

- Suppose we use 3-input LUTs
(implements any function of ≤ 3 inputs, and exactly 1 output)
 - each island must have 1 output and ≤ 3 inputs
 - each circle must be in at least 1 island
 - islands (LUTs) labeled ①, ②, ③, ④..

Place and Route

Synthesis generates netlist -- list of devices and how they're interconnected

LUTs

Place and route determines how to put those devices on a chip and how to lay out wires that connect them

70%

LUTs

Results not as good as you'd like -- ~~40–60%~~ utilization of devices and wires is typical for FGPA

- Can trade off run time of tool for greater utilization to some degree, but there are serious limits
- Beyond 80% utilization, there is a good chance that routing will fail.

Bitstream Generation

- A bitstream in FPGA-speak is a sequence of bits, which
 - Determines how the FPGA fabric is customized in order to implement the design.
 - It determines how IOs, CLBs, and wiring are configured.
- This bitstream is loaded (serially) into the FPGA in a final step.
Now the FPGA is customized to implement the design we wanted.
 - This loading typically takes a few seconds at most.
- Reprogramming simply means that we load a new bitstream on to the FPGA.